

Commodity Machine Learning and beyond

Andreas Müller
NYU, scikit-learn

✉ amueller@nyu.edu

🐦 [@amuellerm1](https://twitter.com/amuellerm1)

🐱 [@amueller](https://github.com/amueller)

🌐 Amueller.io



We Want You!



To Apply Machine Learning!

What ML can do for you

Classification

Hi Andy,

I just received an email from the first tutorial speaker, presenting right before you, saying he's ill and won't be able to make it.

I know you have already committed yourself to two presentations, but is there anyway you could increase your tutorial time slot, maybe just offer time to try out what you've taught? Otherwise I have to do some kind of modern dance interpretation of Python in data :-)
-Leah

Hi Andreas,

I am very interested in your Machine Learning background. I work for X Recruiting who have been engaged by Z, a worldwide leading supplier of Y. We are expanding the core engineering team and we are looking for really passionate engineers who want to create their own story and help millions of people.

Can we find a time for a call to chat for a few minutes about this?

Thanks

Classification

Hi Andy,

I just received an email from the first tutorial speaker, presenting himself as a doctor, saying he's ill and won't be attending.

I know you have a lot to offer yourself to two presentations. I think it would be great if you could increase your audience. Maybe you could just offer time to try out what you've taught? Otherwise I have to do some kind of modern dance interpretation of Python in data :-)
-Leah



Hi Andreas,

I am very interested in your Machine Learning background. I work for X Recruiting who have been engaged by Z, a worldwide leading supplier of Y. We are expanding the core engineering team and we are looking for really passionate engineers who can tell their own story and help millions of people.

Can we find a time to talk for a few minutes about this?

Thanks



Classification

SPONSORED 

[Create Ad](#)

Online Essen bestellen



lieferheld.de



Neu mit Lieferheld: PLZ eingeben, Liefersdienst finden und genießen!

Globaler Chauffeurservice


blacklane.com



Fahren Sie eine Klasse besser zu günstigen Preisen – Blacklane ist weltweit verfügbar!

97,974 people like this

Recommendations



Machine Learning: A Probabilistic Perspective (Adaptive Computation and Machine Learning series) Hardcover

by Kevin P. Murphy (Author)

★★★★★ 31 customer reviews

See all 2 formats and editions

Kindle
\$64.06

Read with our free app

Hardcover
\$77.41

20 Used from \$67.99

42 New from \$69.66

Today's Web-enabled deluge of electronic data calls for automated methods of data analysis. Machine learning provides these, developing methods that can automatically detect patterns in data and then use the uncovered patterns to predict future data. This textbook offers a comprehensive and self-contained introduction to the field of machine learning, based on a unified, probabilistic approach. The coverage combines breadth and depth, offering necessary background material on such topics as probability, optimization, and linear algebra as well as discussion of recent developments in the field, including conditional random fields, L1 regularization, and deep learning. The book is written in an informal, accessible style, complete with pseudo-code for the most important algorithms. All topics are copiously illustrated with color images and worked examples drawn from such application domains as

Read more



See all 1 image(s)

Share

Buy New \$77.41

Qty: 1

List Price: \$99.99

Save: \$12.59 (14%)

FREE Shipping

In Stock.

Ships from and sold by Amazon.com.
Gift-wrap available.



Add to Cart

Sign in to turn on 1-click ordering

Want it Wednesday, June 11? Order within 11 hrs 27 mins and choose AmazonGlobal Priority Shipping at checkout. [Details](#)

Buy Used \$71.98

Add to Wish List

Trade in your item
Get a \$35.10
Gift Card.

Trade in

[Learn More](#)

Have one to sell?

[Sell on Amazon](#)

Frequently Bought Together



Price for all three: \$249.90

Add all three to Cart

Add all three to Wish List

Show availability and shipping details

✓ **This Item:** Machine Learning: A Probabilistic Perspective (Adaptive Computation and Machine Learning series) by Kevin P. Murphy Hardcover \$77.41

✓ Pattern Recognition and Machine Learning (Information Science and Statistics) by Christopher M. Bishop Hardcover \$85.46

✓ Probabilistic Graphical Models: Principles and Techniques (Adaptive Computation and Machine Learning ... by Daphne Koller Hardcover \$87.03

Customers Who Bought This Item Also Bought

Page 1 of 12



Learning From Data
Yaser S. Abu-Mostafa



Pattern Recognition and
Machine Learning ...



Probabilistic Graphical
Models: Principles ...



The Elements of Statistical
Learning: ...



Foundations of Machine
Learning (Adaptive ...



Machine Learning: The Art
and Science of ...



Building Machine Learning
Systems with ...

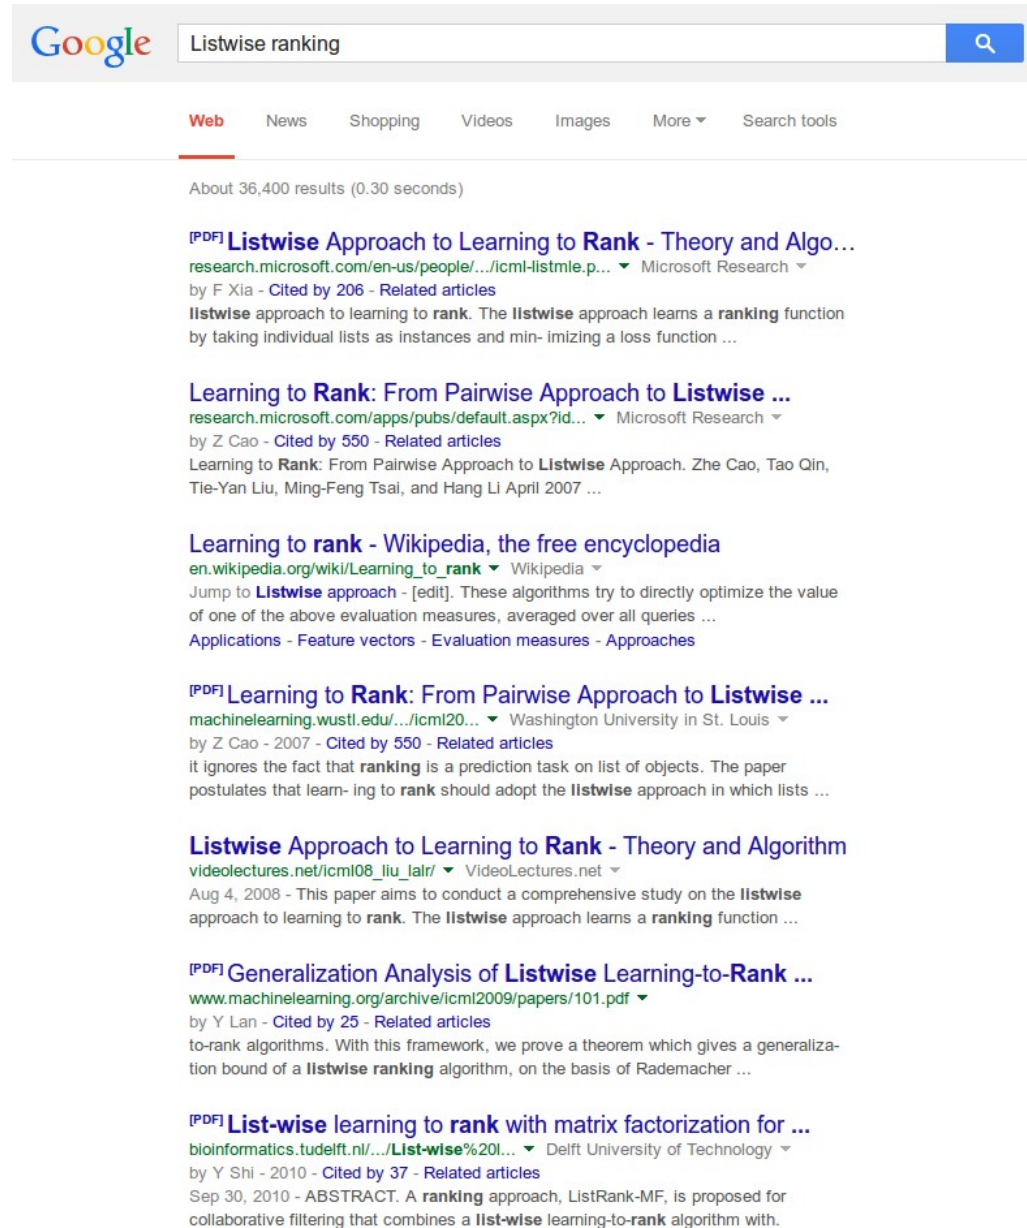


Understanding Machine
Learning: From Theory to ...



An Introduction to
Statistical Learning: ...

Ranking



The image shows a Google search results page for the query "Listwise ranking". The search bar at the top contains the text "Listwise ranking" and a magnifying glass icon. Below the search bar, there are tabs for "Web", "News", "Shopping", "Videos", "Images", "More", and "Search tools". The "Web" tab is selected. The search results show "About 36,400 results (0.30 seconds)".

The first result is a PDF titled "Listwise Approach to Learning to Rank - Theory and Algo..." from Microsoft Research, by F Xia. It is cited by 206 and has related articles. The abstract states: "Listwise approach to learning to rank. The listwise approach learns a ranking function by taking individual lists as instances and minimizing a loss function ...".

The second result is "Learning to Rank: From Pairwise Approach to Listwise ..." from Microsoft Research, by Z Cao. It is cited by 550 and has related articles. The abstract states: "Learning to Rank: From Pairwise Approach to Listwise Approach. Zhe Cao, Tao Qin, Tie-Yan Liu, Ming-Feng Tsai, and Hang Li April 2007 ...".

The third result is "Learning to rank - Wikipedia, the free encyclopedia". It includes a link to the Wikipedia page and a summary: "Jump to Listwise approach - [edit]. These algorithms try to directly optimize the value of one of the above evaluation measures, averaged over all queries ...". It also lists "Applications - Feature vectors - Evaluation measures - Approaches".

The fourth result is a PDF titled "Learning to Rank: From Pairwise Approach to Listwise ..." from Washington University in St. Louis, by Z Cao. It is cited by 550 and has related articles. The abstract states: "it ignores the fact that ranking is a prediction task on list of objects. The paper postulates that learning to rank should adopt the listwise approach in which lists ...".

The fifth result is "Listwise Approach to Learning to Rank - Theory and Algorithm" from VideoLectures.net, dated Aug 4, 2008. The abstract states: "This paper aims to conduct a comprehensive study on the listwise approach to learning to rank. The listwise approach learns a ranking function ...".

The sixth result is a PDF titled "Generalization Analysis of Listwise Learning-to-Rank ..." from machinelearning.org, by Y Lan. It is cited by 25 and has related articles. The abstract states: "to-rank algorithms. With this framework, we prove a theorem which gives a generalization bound of a listwise ranking algorithm, on the basis of Rademacher ...".

The seventh result is a PDF titled "List-wise learning to rank with matrix factorization for ..." from Delft University of Technology, by Y Shi. It is cited by 37 and has related articles. The abstract states: "Sep 30, 2010 - ABSTRACT. A ranking approach, ListRank-MF, is proposed for collaborative filtering that combines a list-wise learning-to-rank algorithm with."

Applying machine learning is easy.

Applying machine learning is easy.
But it should be easier!




python


IP[y]: IPython
Interactive Computing



SciPy






NumPy



matplotlib

pandas

$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$



IP[y]: IPython Interactive Computing

```
Chapter 4 - Working With Text Data. - Mozilla Firefox
File Edit View History Bookmarks Tools Help
IPy Chapter 4 - Workin...
127.0.0.1:8888/030923ca-af46-4fc6-86fc-1f0b18130e27#
DuckDuckGo
IP[y]: Notebook Chapter 4 - Working With Text Data. (autosaved)
File Edit View Insert Cell Kernel Help
Code Cell Toolbar: None

In [22]: from sklearn.svm import LinearSVC
svm = LinearSVC(C=0.01)

In [23]: svm.fit(X_train, y_train)

Out[23]: LinearSVC(C=0.01, class_weight=None, dual=True, fit_intercept=True,
intercept_scaling=1, loss='l2', multi_class='ovr', penalty='l2',
random_state=None, tol=0.0001, verbose=0)

In [24]: svm.score(X_train, y_train)

Out[24]: 0.88421586014694709

In [25]: svm.score(X_test, y_test)

Out[25]: 0.83679637325273892

In [26]: y_test_pred = svm.predict(X_test)

In [27]: from sklearn.metrics import classification_report


In [28]: print(classification_report(y_test, y_test_pred))

              precision    recall  f1-score   support

     0       0.84         0.96      0.90       1954
     1       0.82         0.48      0.61        693




 avg / total       0.83         0.84      0.82       2647

In [29]: coef = svm.coef_.ravel()
positive_coefficients = np.argsort(coef)[-25:]
negative_coefficients = np.argsort(coef)[:25]
interesting_coefficients = np.hstack([negative_coefficients, positive_coefficients])
```



pandas

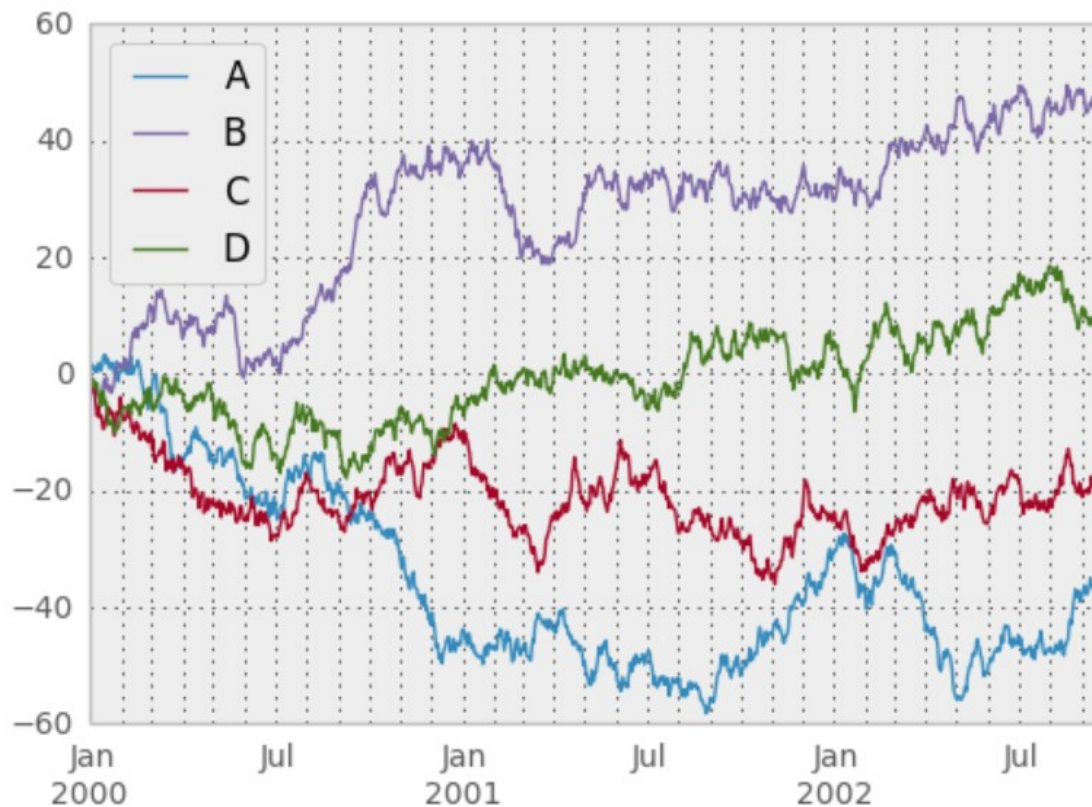
$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$

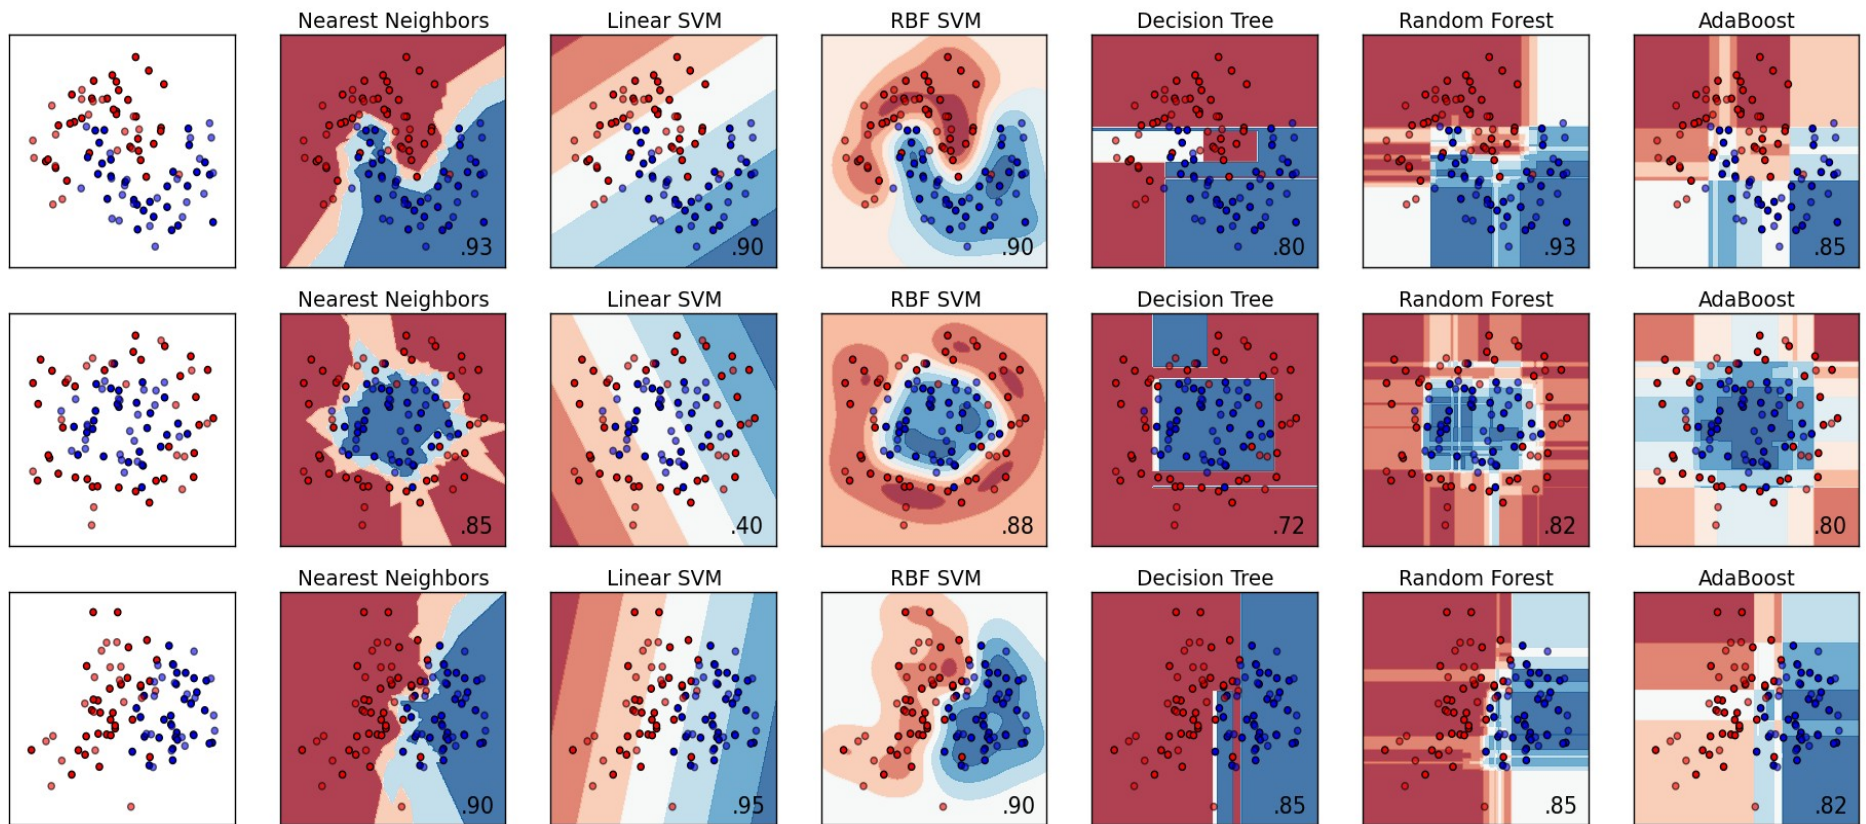


```
In [6]: df = DataFrame(randn(1000, 4), index=ts.index, columns=list('ABCD'))
```

```
In [7]: df = df.cumsum()
```

```
In [8]: plt.figure(); df.plot();
```





lovely

The
New York
Times

Booking.com



2750+ research papers

Fully Functional Spam Classifier

```
from sklearn.naive_bayes import MultinomialNB
from sklearn.feature_extraction.text import CountVectorizer
from pipeline import make_pipeline

spam_classifier = make_pipeline(CountVectorizer(),
                                MultinomialNB())
spam_classifier.fit(email_texts, is_spam)

spam_classifier.predict(new_emails)
```

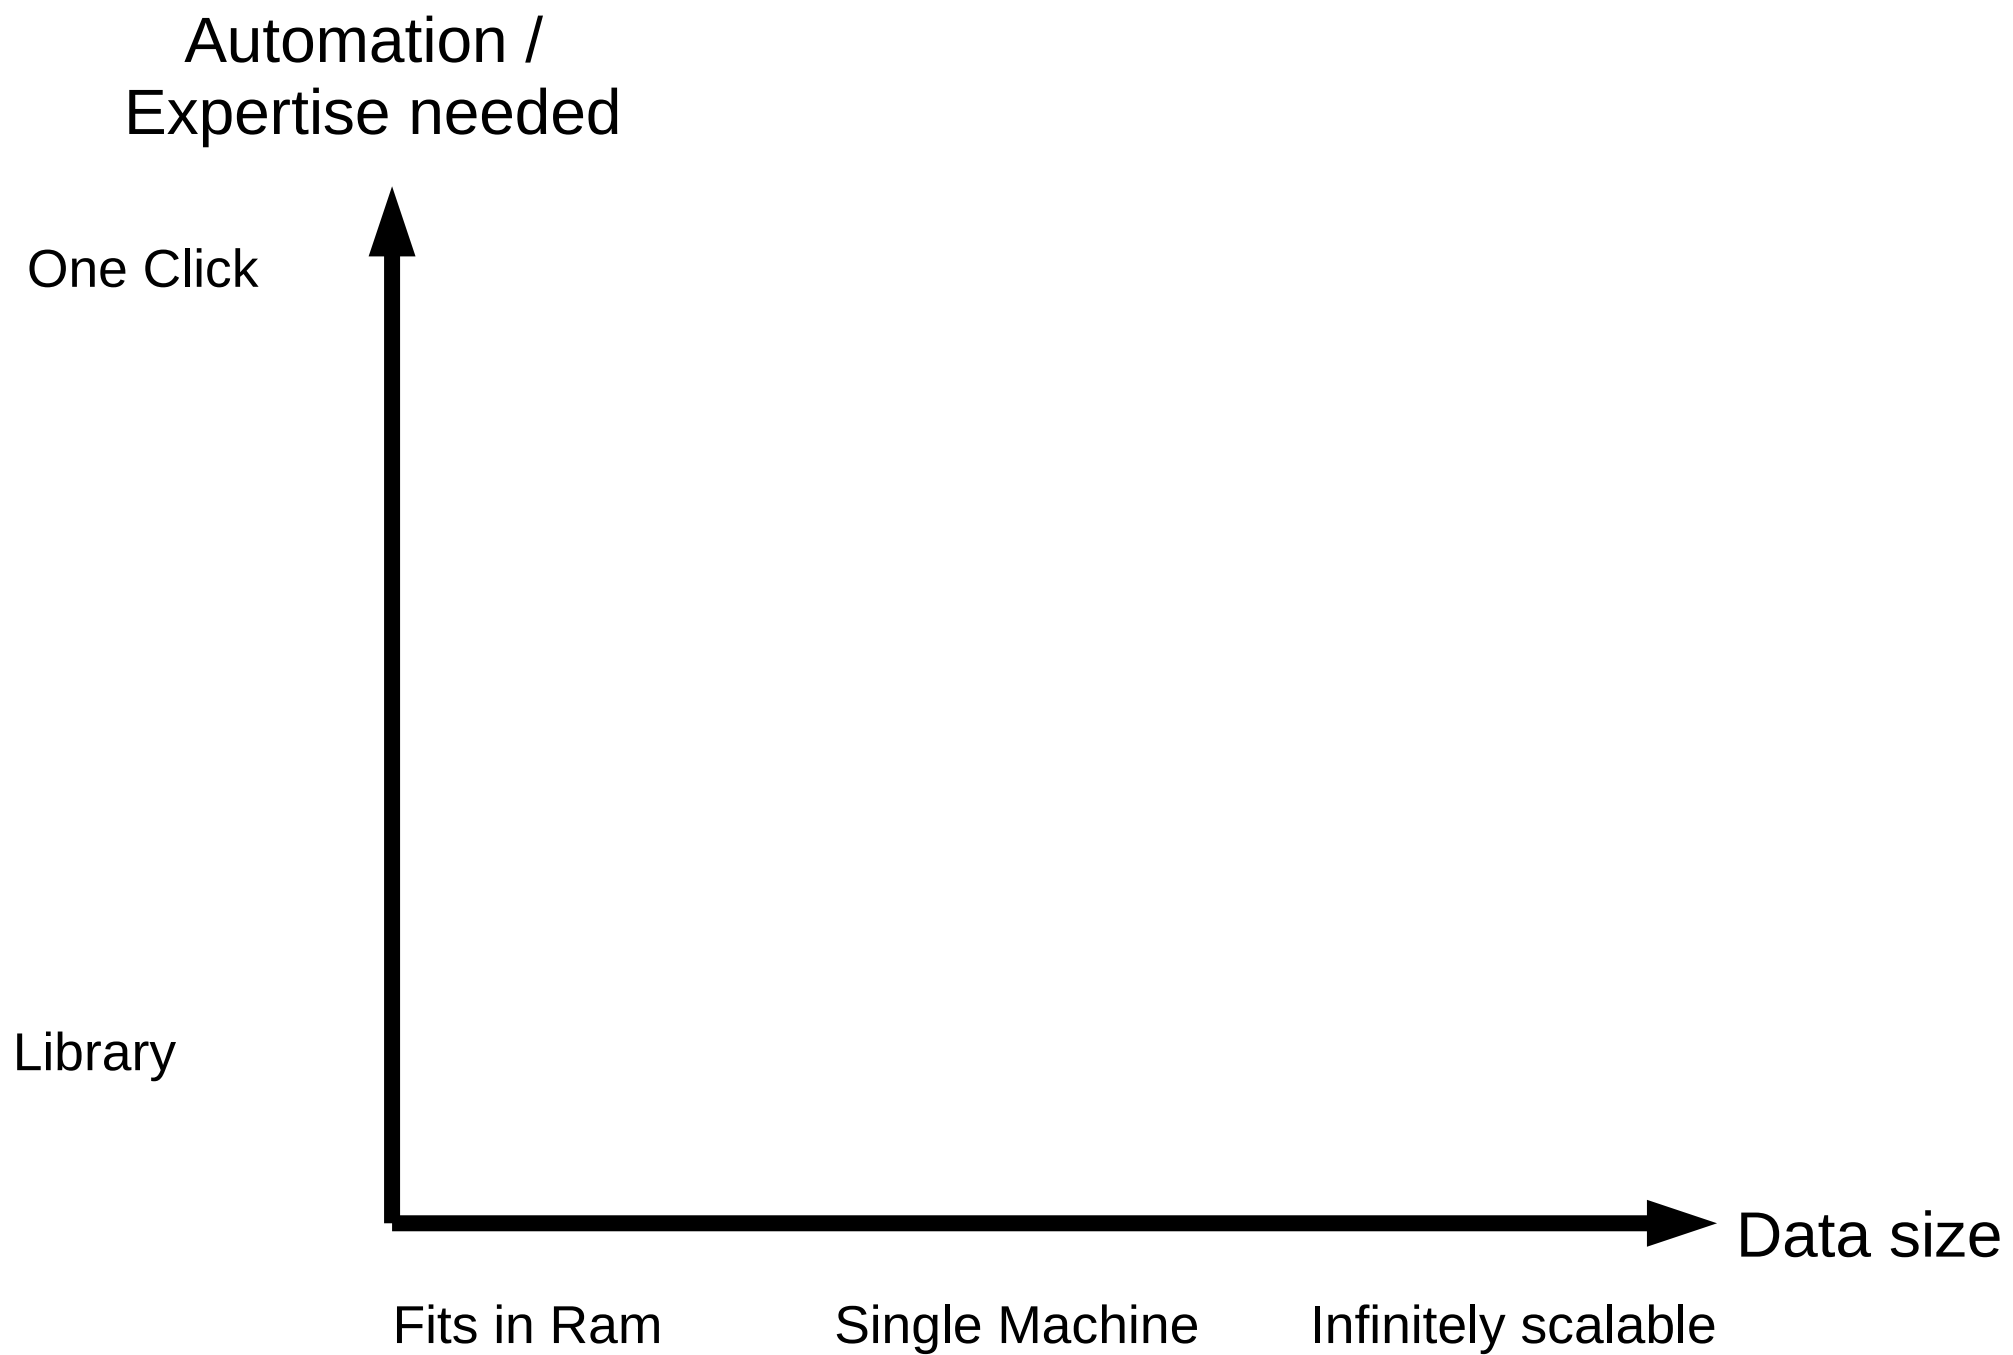
“The scikit-learn tutorials / documentation is so good, one doesn't need a textbook anymore to learn a new machine learning method.”

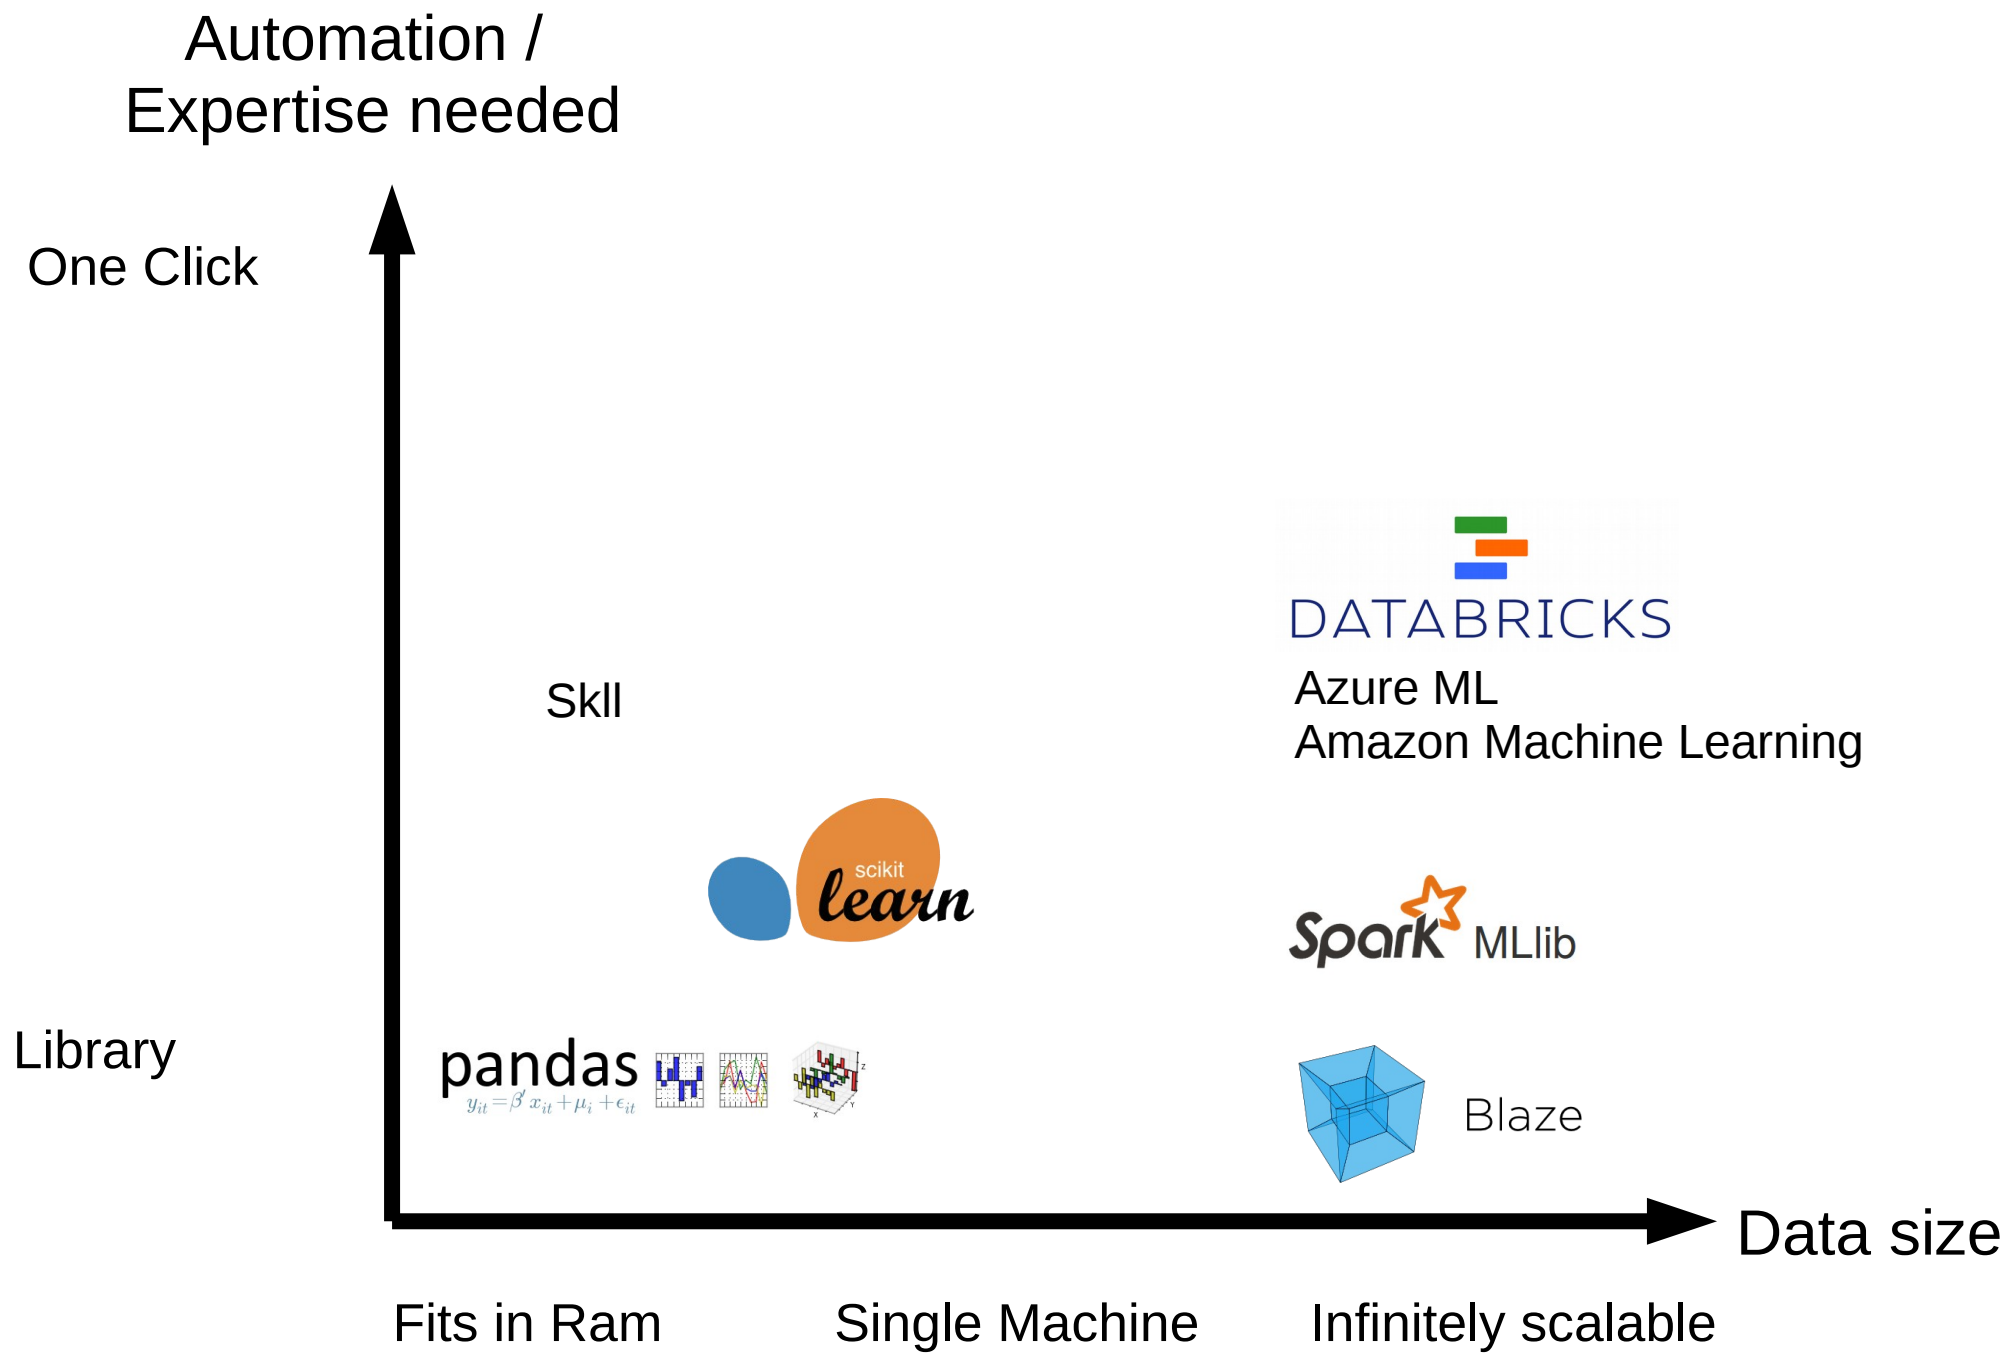
This is not enough!

Automation /
Expertise needed



Data size





A single machine is (usually) enough.

	vCPU	ECU	Memory (GiB)	Instance Storage (GB)	Linux/UNIX Usage
r3.large	2	6.5	15	1 x 32 SSD	\$0.166 per Hour
r3.xlarge	4	13	30.5	1 x 80 SSD	\$0.333 per Hour
r3.2xlarge	8	26	61	1 x 160 SSD	\$0.665 per Hour
r3.4xlarge	16	52	122	1 x 320 SSD	\$1.33 per Hour
r3.8xlarge	32	104	244	2 x 320 SSD	\$2.66 per Hour

Storage Optimized - Current Generation

i2.xlarge	4	14	30.5	1 x 800 SSD	\$0.853 per Hour
i2.2xlarge	8	27	61	2 x 800 SSD	\$1.705 per Hour
i2.4xlarge	16	53	122	4 x 800 SSD	\$3.41 per Hour
i2.8xlarge	32	104	244	8 x 800 SSD	\$6.82 per Hour
d2.xlarge	4	14	30.5	3 x 2000 HDD	\$0.69 per Hour
d2.2xlarge	8	28	61	6 x 2000 HDD	\$1.38 per Hour
d2.4xlarge	16	56	122	12 x 2000 HDD	\$2.76 per Hour
d2.8xlarge	36	116	244	24 x 2000 HDD	\$5.52 per Hour

We need open box methods.

We need black-box methods.

MLService

predict

Hyperparameter Optimization


Spearmint

Hyperopt

SMAC

... scikit-learn

Bayesian Optimization

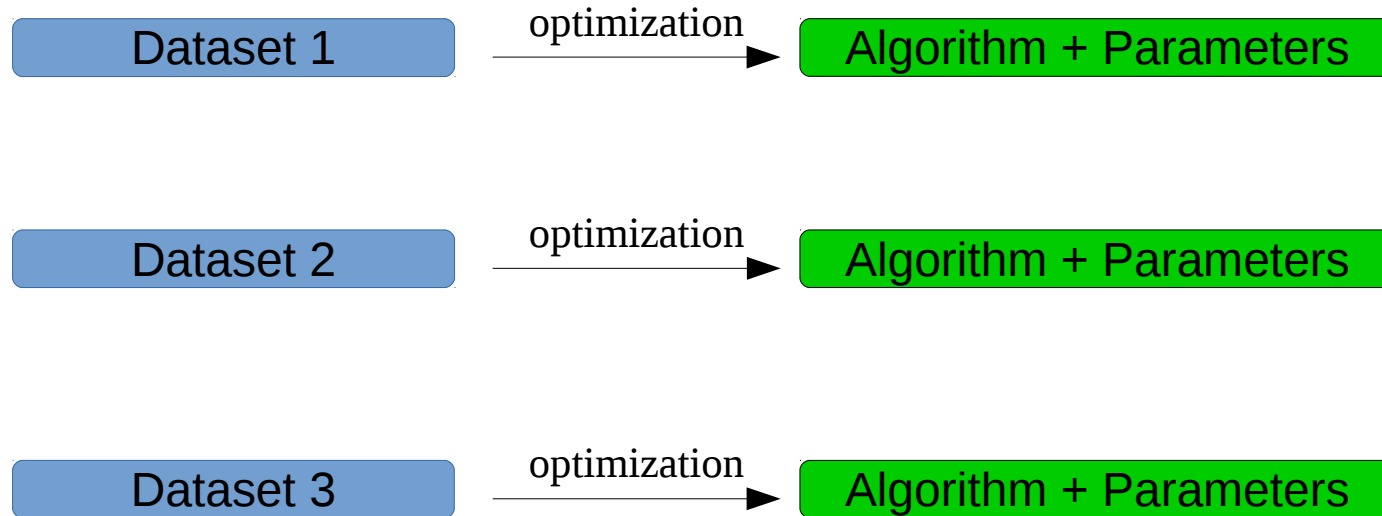


From Eric Brochu, Vlad M. Cora and Nando de Freitas

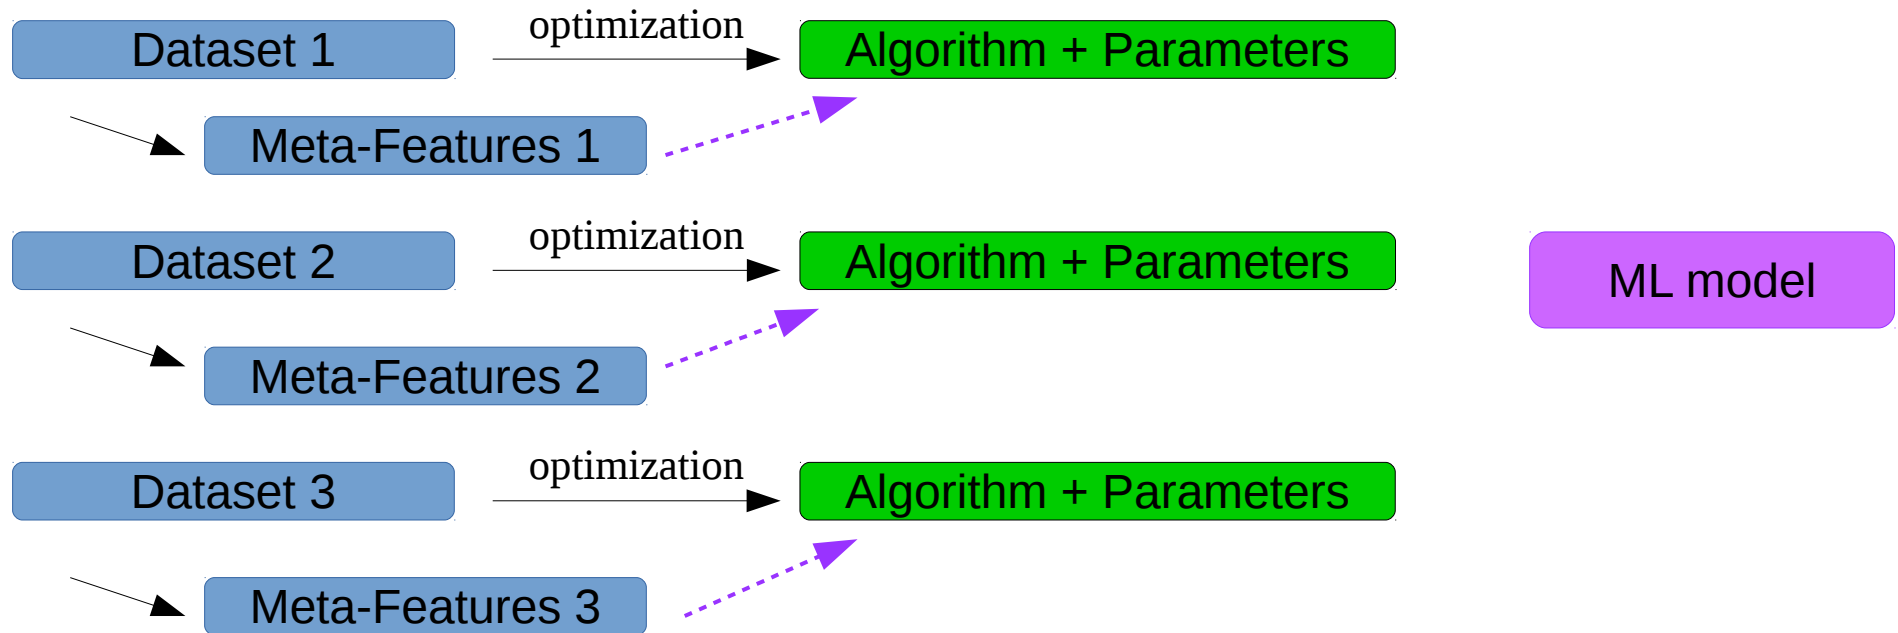
Meta-Learning



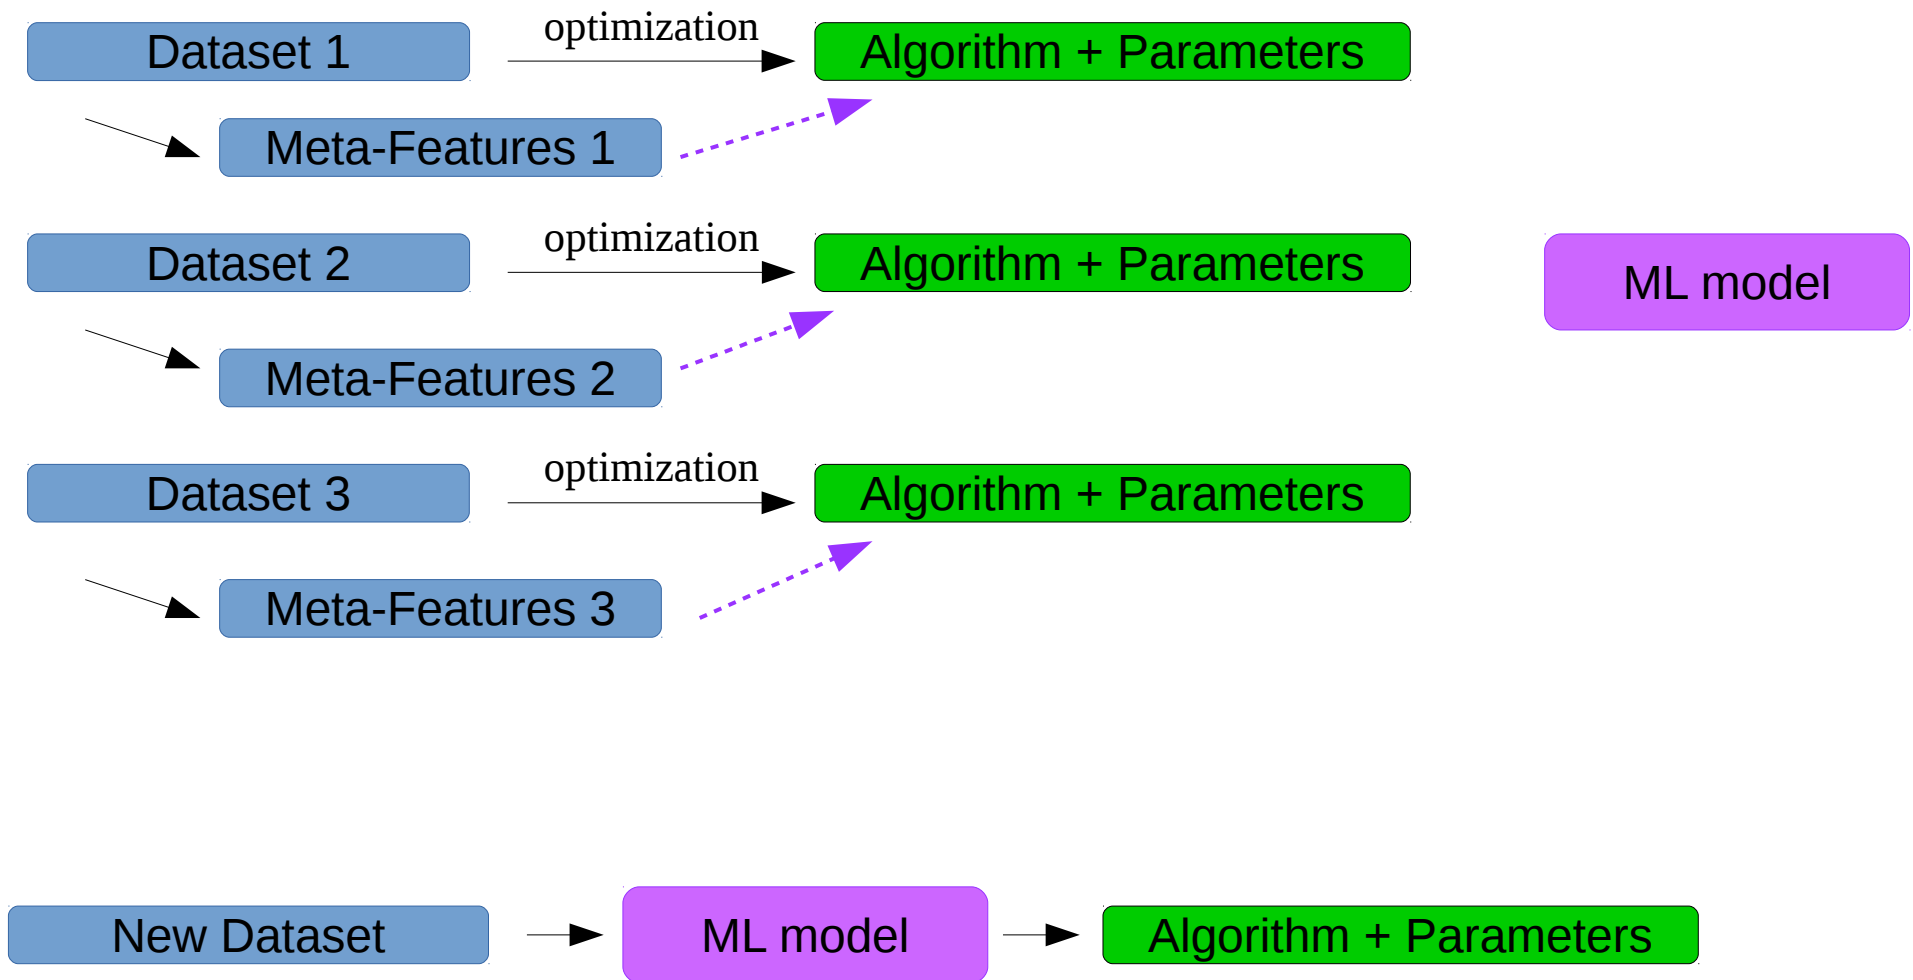
Meta-Learning




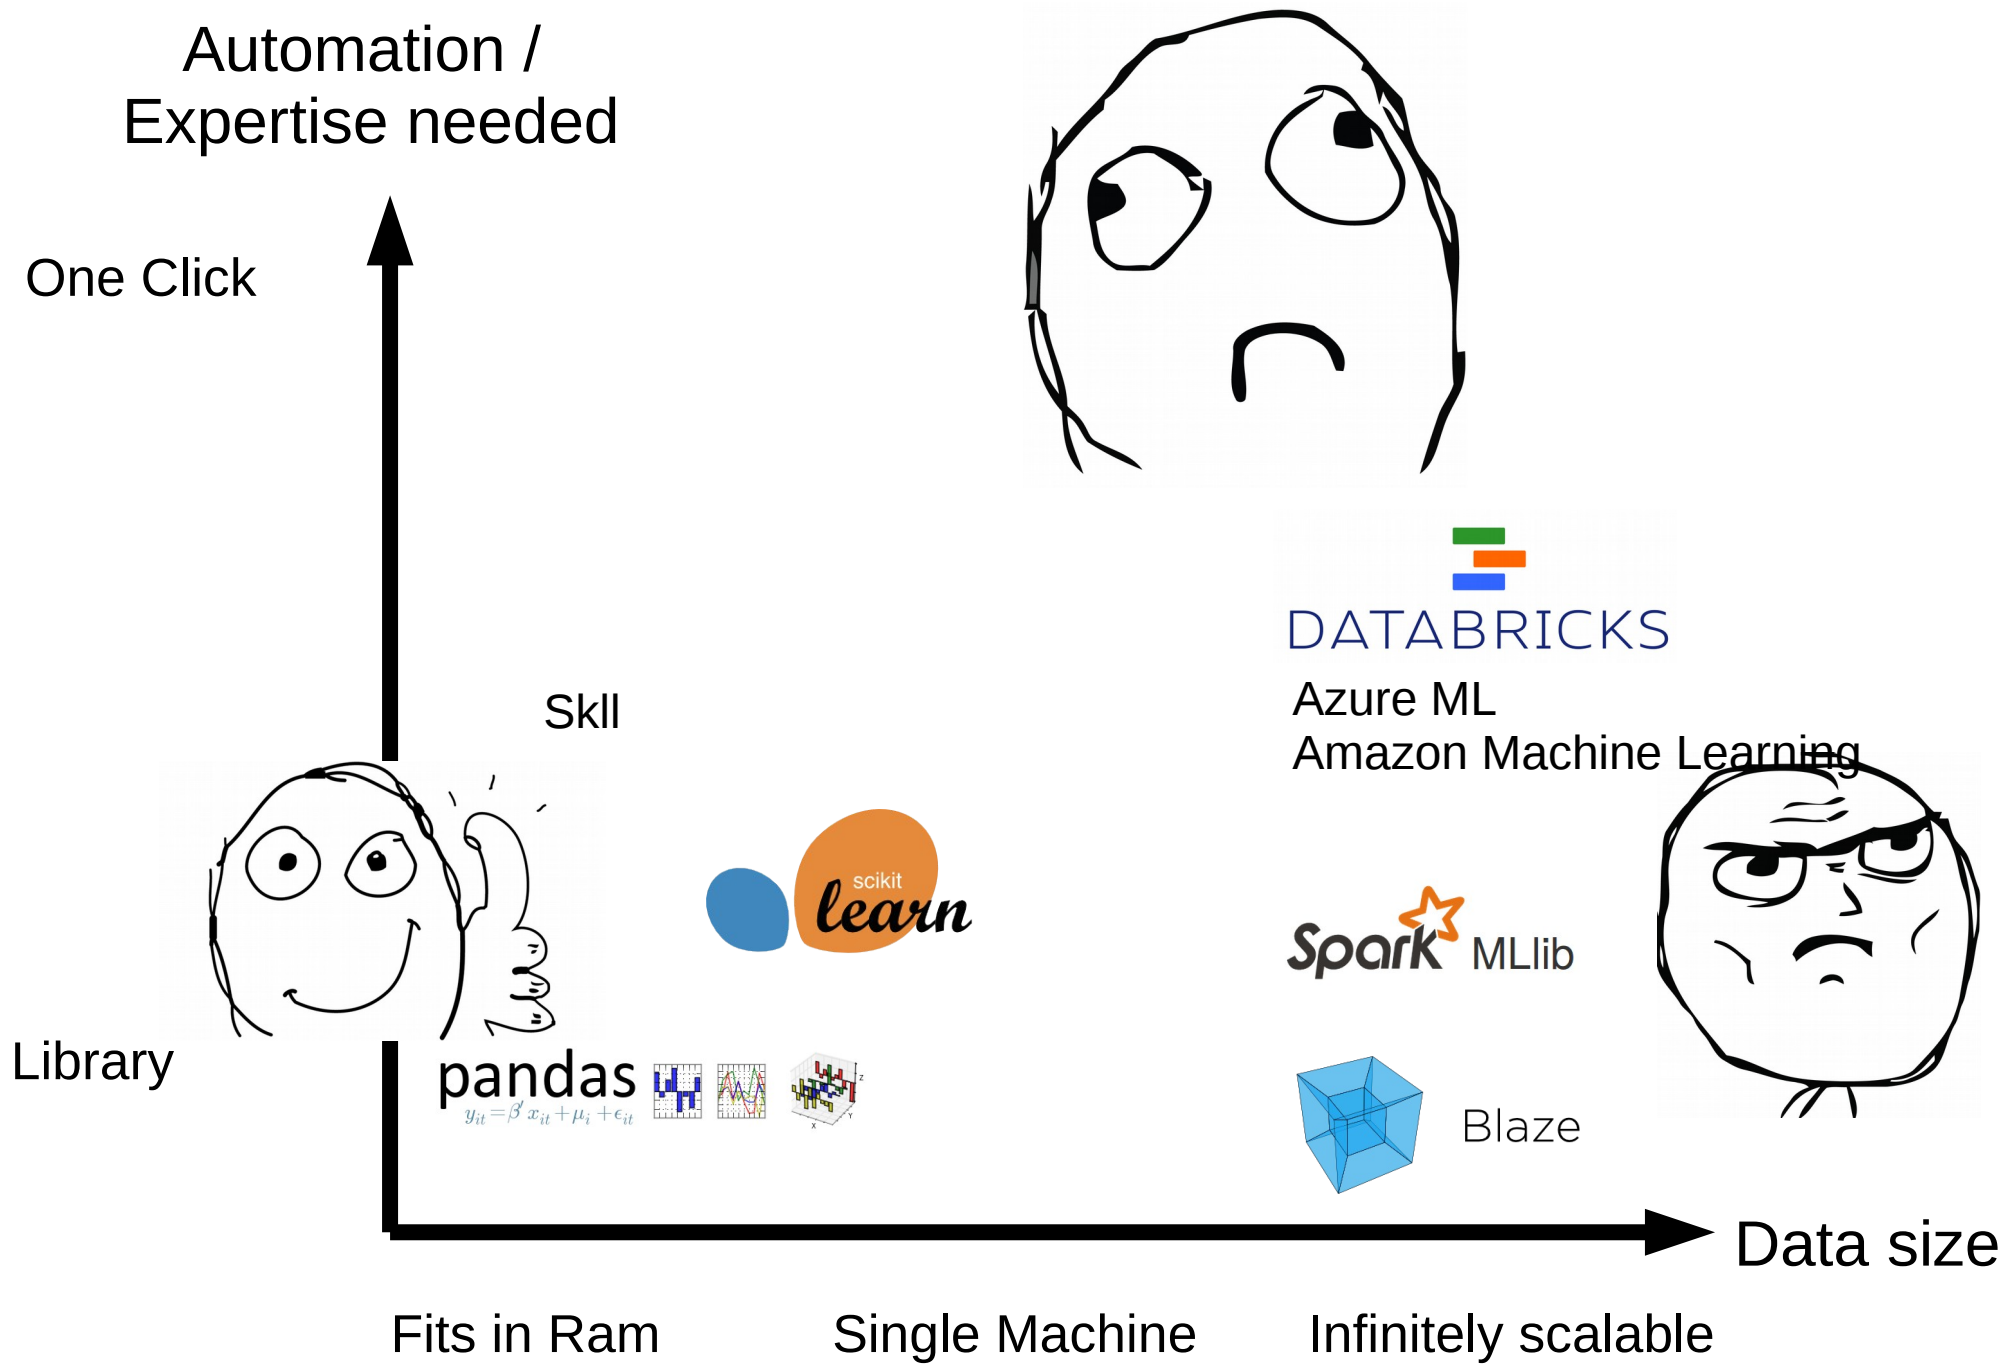
Meta-Learning

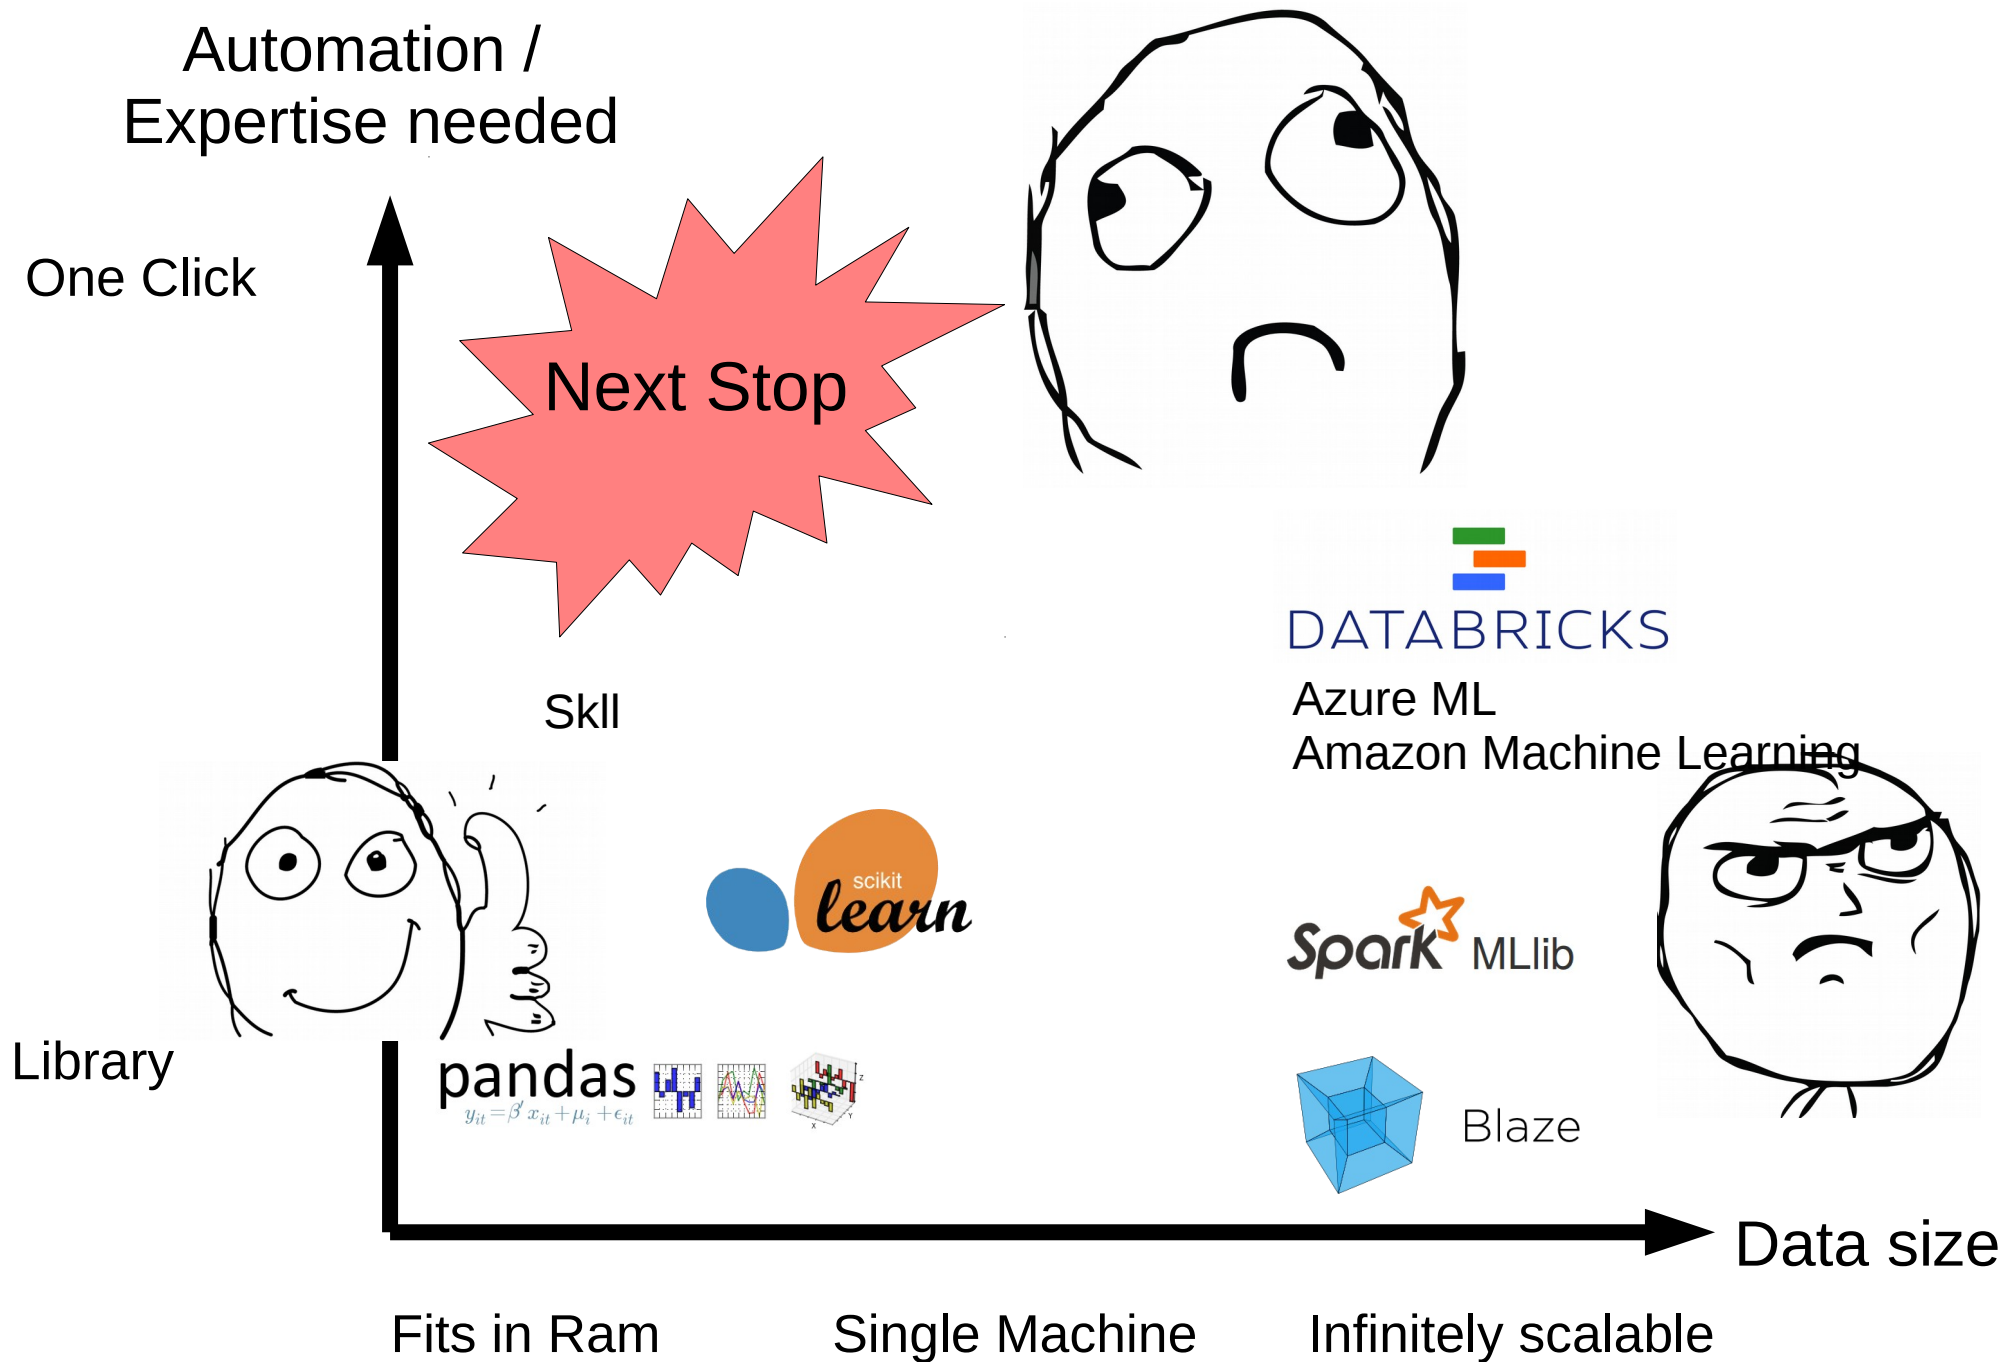


Meta-Learning









O'REILLY®



Introduction to Machine Learning with Python

A GUIDE FOR DATA SCIENTISTS

Andreas C. Müller & Sarah Guido

Release June 2016

Thank you.



@amuellermi



@amueller



amueller@nyu.edu



<http://amueller.io>