

# Automatic Machine Learning?

Andreas Müller

Columbia University, scikit-learn



Alfred P. Sloan  
FOUNDATION



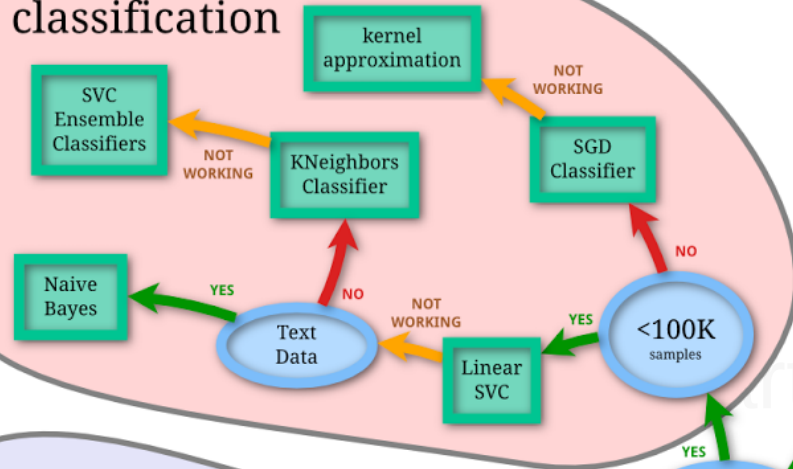
COLUMBIA UNIVERSITY  
IN THE CITY OF NEW YORK

Why?

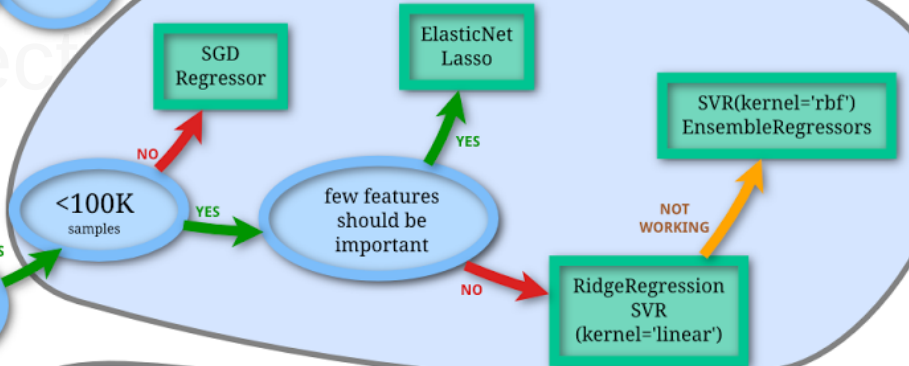
Issues with current tools (scikit-learn)

# scikit-learn algorithm cheat-sheet

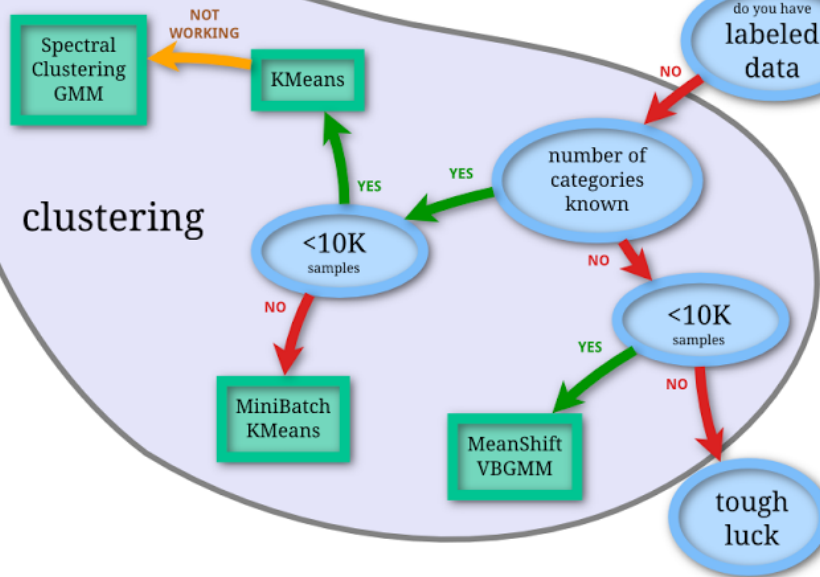
## classification



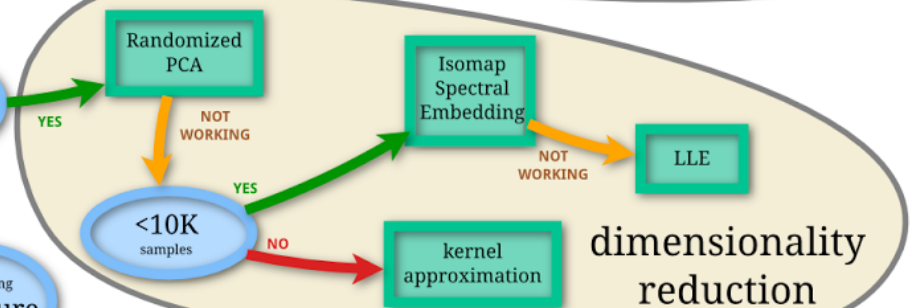
## regression



## clustering



## dimensionality reduction



# Selecting Hyper-Parameters

```
In [2]: clf = SVC()  
        clf.fit(X_train, y_train)
```

```
SVC(self, C=1.0, kernel='rbf', degree=3, gamma=0.0, coef0=0.0,  
     shrinking=True, probability=False, tol=0.001, cache_size=200,  
     class_weight=None, verbose=False, max_iter=-1, random_state=None)
```

# Scikit-learn: Explicit is better than implicit

```
make_pipeline(  
    OneHotEncoder(),  
    Imputer(),  
    StandardScaler(),  
    SVC())
```

# What?

```
from automl import AutoClassifier  
clf = AutoClassifier().fit(X_train, y_train)
```

```
> Current Accuracy: 70% (AUC .65) LinearSVC(C=1), 10sec  
> Current Accuracy: 76% (AUC .71) RandomForest(n_estimators=20) 30sec  
> Current Accuracy: 80% (AUC .74) RandomForest(n_estimators=500) 30sec
```

## Step 1: Automate Parameter Selection



## Step 2: Automate Model Selection

## Step 3: Automate Pipeline Selection

How?

# Formalizing the Search Space

Discrete and Continuous Parameters

Conditional Parameters

Fixed pipeline vs flexible pipeline

# Formalizing the Search Space

Discrete and Continuous Parameters

Conditional Parameters

Fixed pipeline vs flexible pipeline

# Search Methods

# Exhaustive Search (Grid Search)

# Randomized Search




# Randomized Search

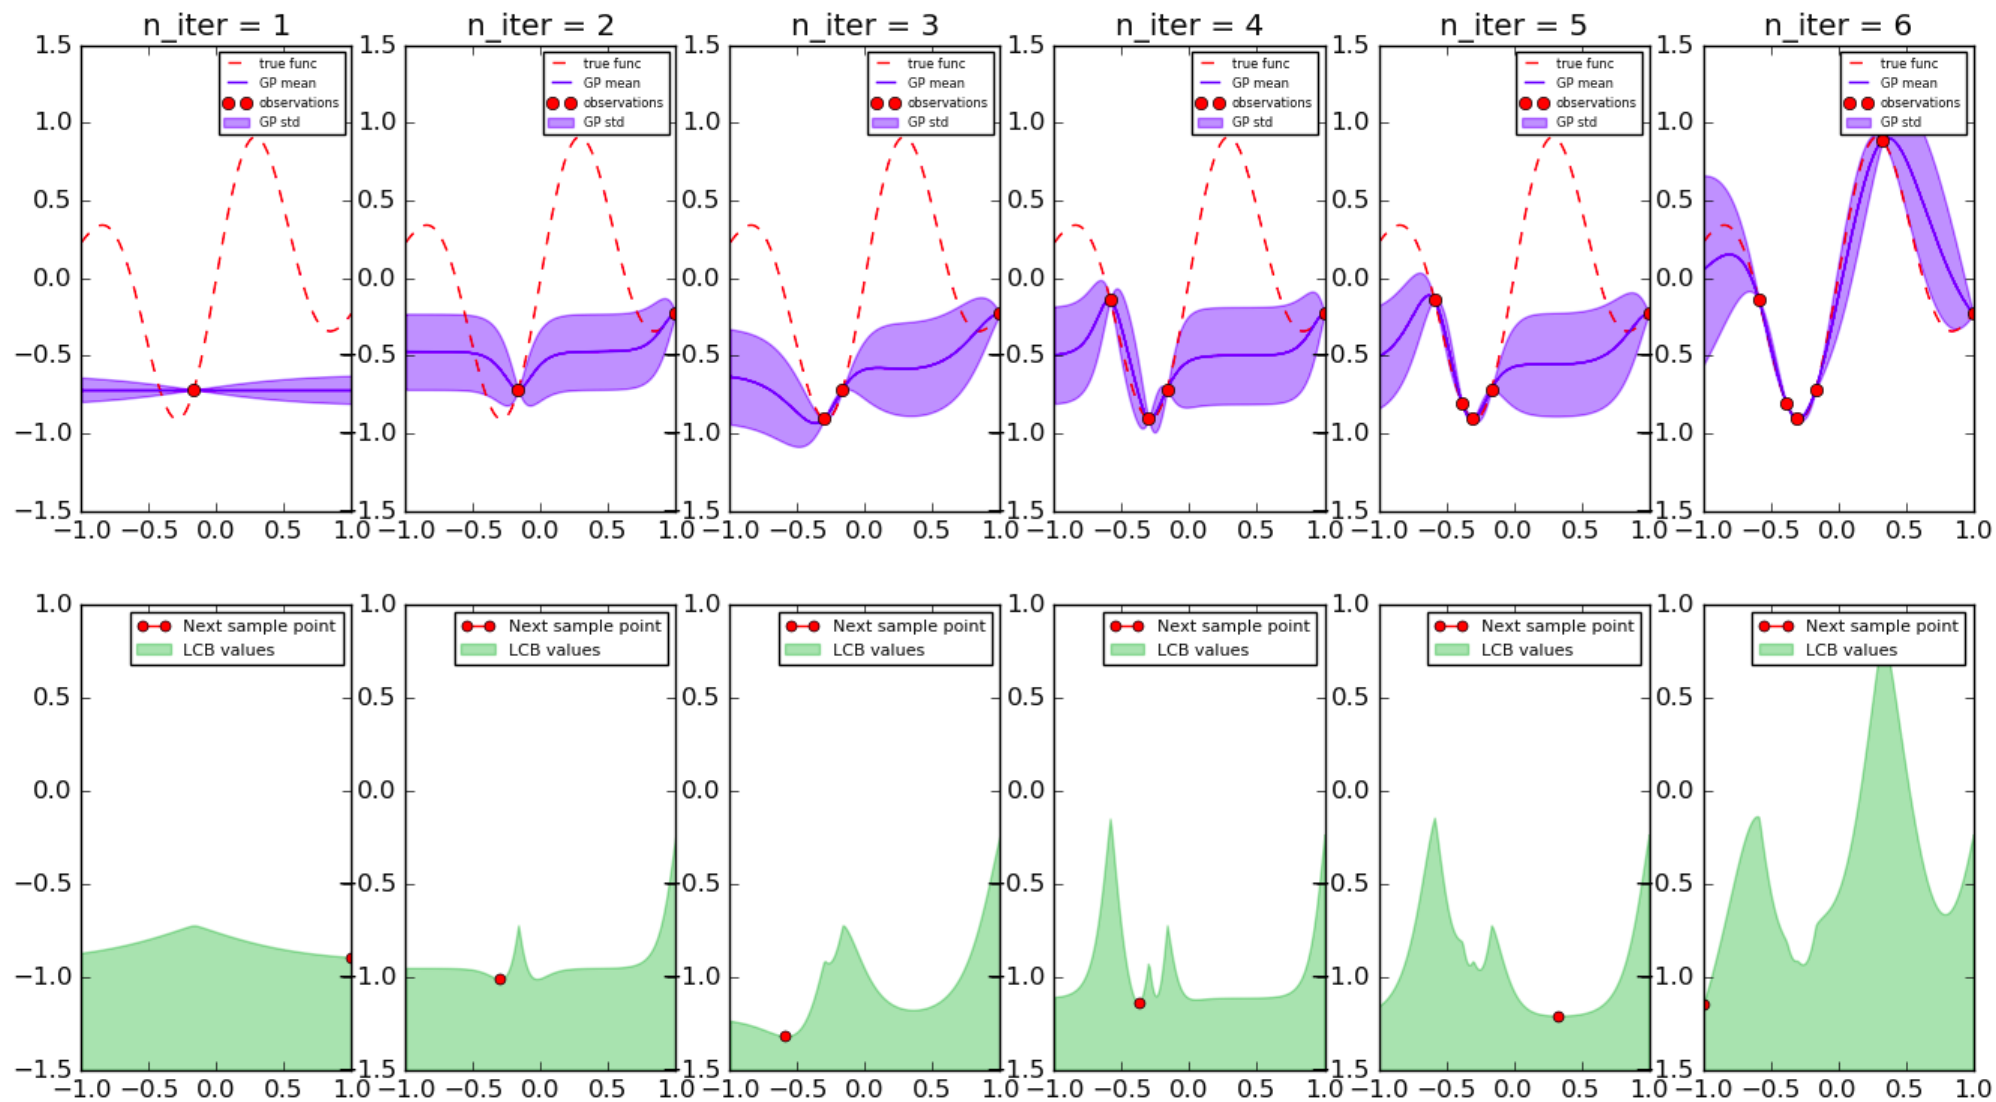
# Bayesian Optimization (SMBO)

# Bayesian Optimization (SMBO)


# Gaussian process based minimization



# Gaussian process based minimization



# Gaussian process based minimization



# Common Models

|                             | Discrete<br>Parameters | Scalable | Conditional<br>Parameters | Papers &<br>implementations | Specify<br>parameters |
|-----------------------------|------------------------|----------|---------------------------|-----------------------------|-----------------------|
| Gaussian<br>Process         | ?                      | ✗        | ?                         | many                        | bounds                |
| Random<br>Forest<br>(SMAC)  | ✓                      | ✓        | ✓                         | few                         | prior                 |
| Non-<br>parametric<br>(TPE) | ✓                      | ?        | ✓                         | few                         | prior                 |

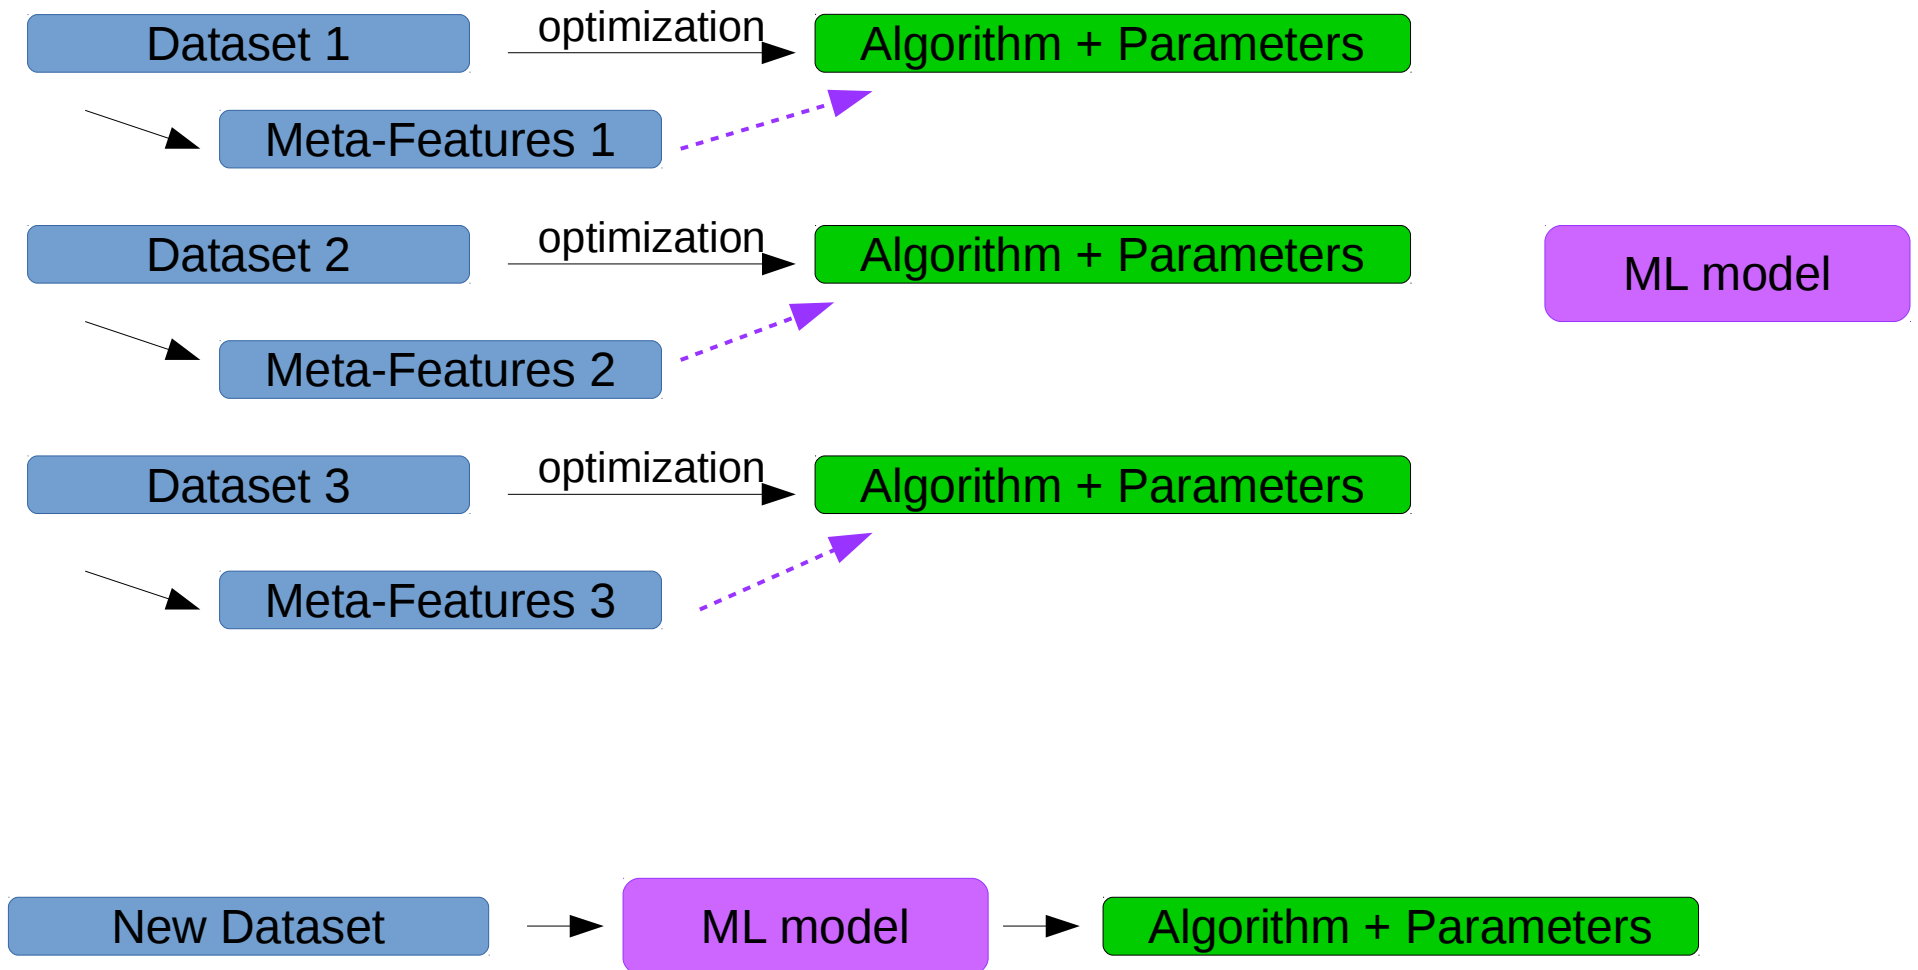
# Common Models

|                             | Discrete<br>Parameters | Scalable | Conditional<br>Parameters | Papers &<br>implementations | Specify<br>parameters |
|-----------------------------|------------------------|----------|---------------------------|-----------------------------|-----------------------|
| Gaussian<br>Process         | ?                      | ✗        | ?                         | many                        | bounds                |
| Random<br>Forest<br>(SMAC)  | ✓                      | ✓        | ✓                         | few                         | prior                 |
| Non-<br>parametric<br>(TPE) | ✓                      | ?        | ✓                         | few                         | prior                 |



# Warm-starting and Meta-learning

# Meta-Learning



# Meta-Features

# Existing Approaches

# auto-sklearn

(Hutter, Feurer, Eggensperger)

<http://automl.github.io/auto-sklearn/stable/>


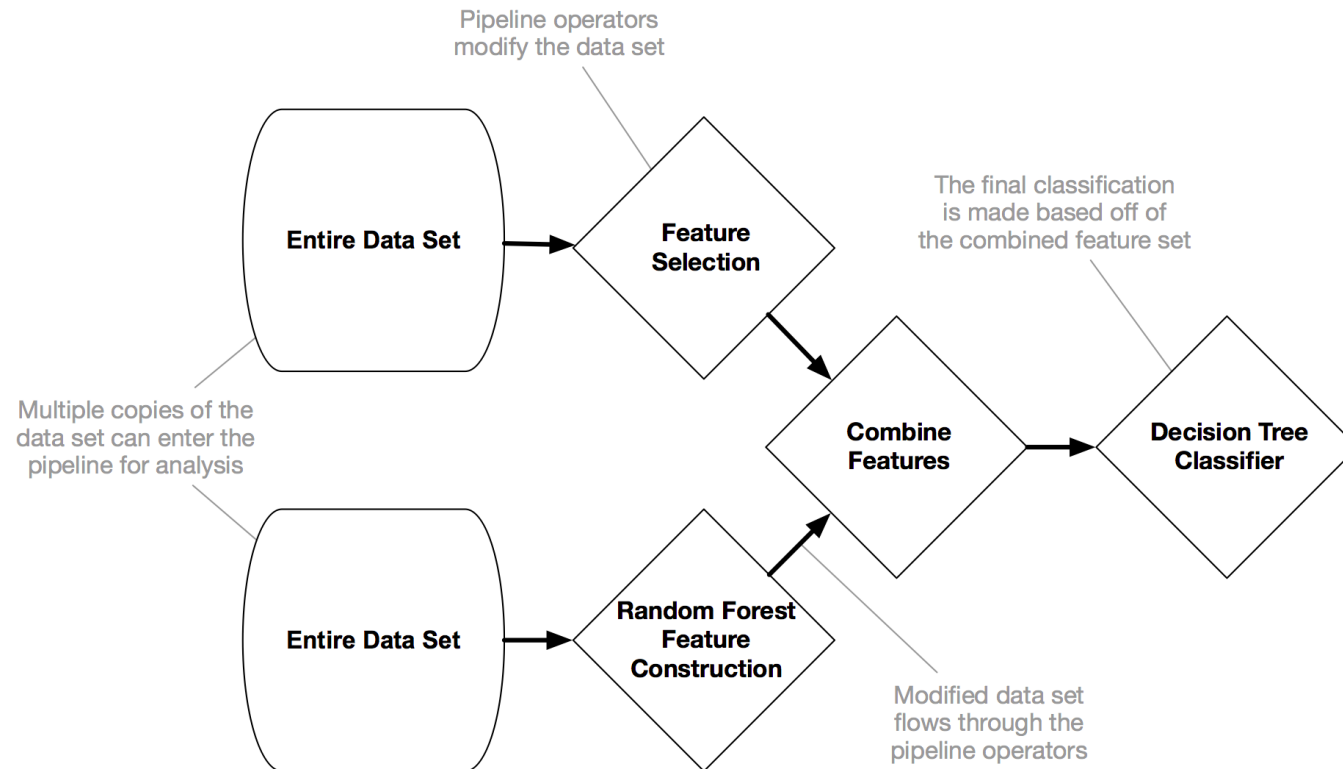
# Autoweka

<http://www.cs.ubc.ca/labs/beta/Projects/autoweka/>

# Hyperopt-sklearn

<http://hyperopt.github.io/hyperopt-sklearn/>

# TPot





# SMBO Packages

# Spearmint

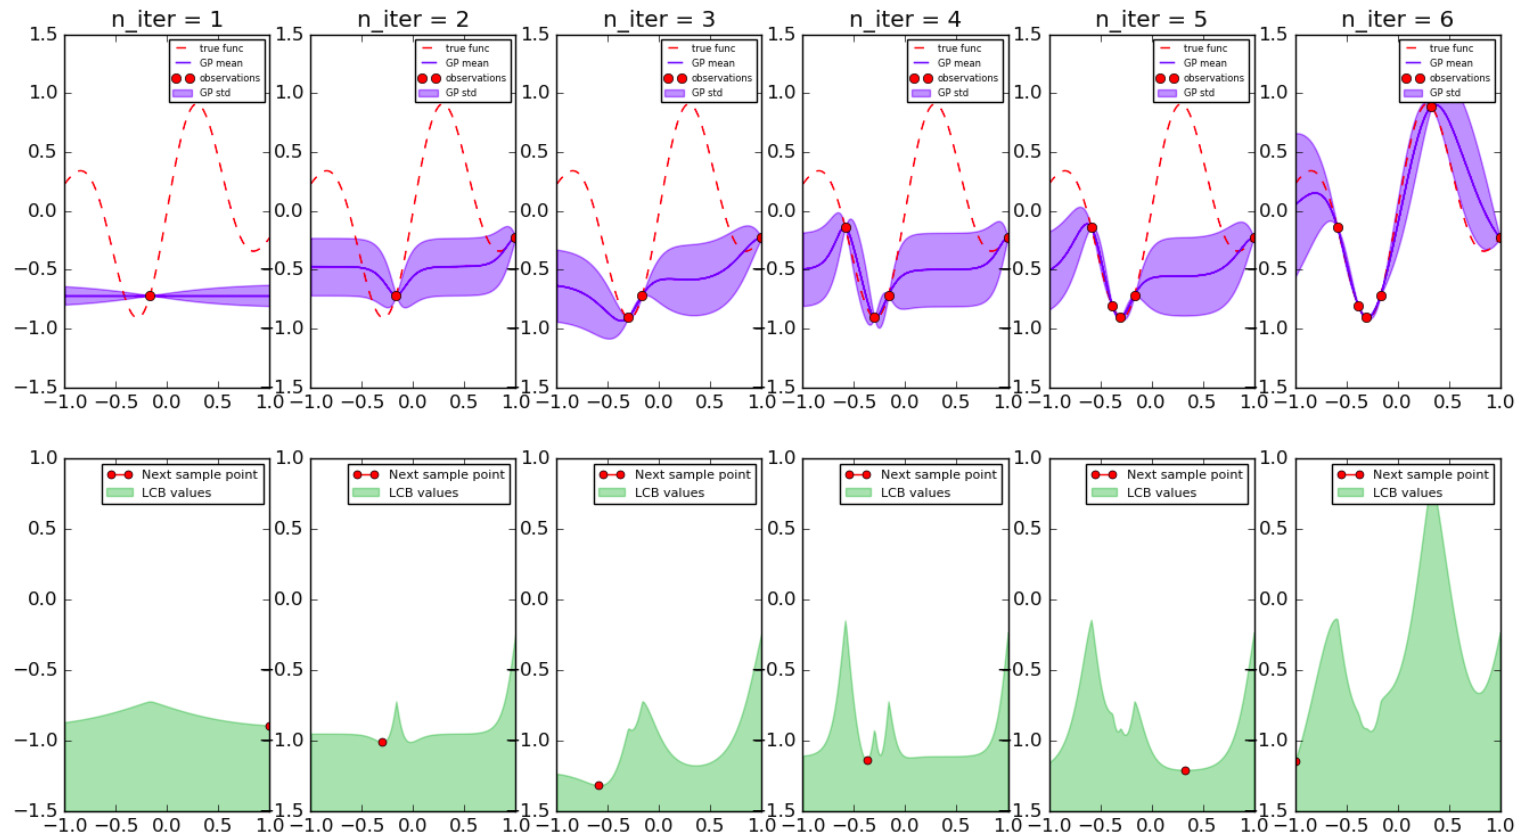
<https://github.com/HIPS/Spearmint>

# GPyOpt

<https://github.com/SheffieldML/GPyOpt>

# Scikit-optimize

Gaussian process based minimization



# Some Benchmarks

| Experiment          | # Evals | SMAC                             | TPE               | Spearmint                            | DNGO                               |
|---------------------|---------|----------------------------------|-------------------|--------------------------------------|------------------------------------|
| Branin (0.398)      | 200     | $0.655 \pm 0.27$                 | $0.526 \pm 0.13$  | <b><math>0.398 \pm 0.00</math></b>   | <b><math>0.398 \pm 0.00</math></b> |
| Hartmann6 (-3.322)  | 200     | $-2.977 \pm 0.11$                | $-2.823 \pm 0.18$ | <b><math>-3.3166 \pm 0.02</math></b> | $-3.319 \pm 0.00$                  |
| Logistic Regression | 100     | $8.6 \pm 0.9$                    | $8.2 \pm 0.6$     | <b><math>6.88 \pm 0.0</math></b>     | <b><math>6.89 \pm 0.04</math></b>  |
| LDA (On grid)       | 50      | $1269.6 \pm 2.9$                 | $1271.5 \pm 3.5$  | <b><math>1266.2 \pm 0.1</math></b>   | <b><math>1266.2 \pm 0.0</math></b> |
| SVM (On grid)       | 100     | <b><math>24.1 \pm 0.1</math></b> | $24.2 \pm 0.0$    | <b><math>24.1 \pm 0.1</math></b>     | <b><math>24.1 \pm 0.1</math></b>   |

Results below are for `n_calls=64` :

| Method                         | Average rank (less is better) |
|--------------------------------|-------------------------------|
| <code>dummy_minimize</code>    | 4.552                         |
| <code>forest_minimize</code>   | 2.362                         |
| <code>gbrt_minimize</code>     | 2.172                         |
| <code>gp_minimize</code>       | 1.241                         |
| <code>gpyopt_minimize</code>   | 1.069                         |
| <code>hyperopt_minimize</code> | 3.052                         |
| <code>smac_minimize</code>     | 3.431                         |


} scikit-optimize

GpyOpt (GP)


Hyperopt (TPE)

SMAC (Random Forest)


<https://github.com/iaroslav-ai/scikit-optimize-benchmarks>



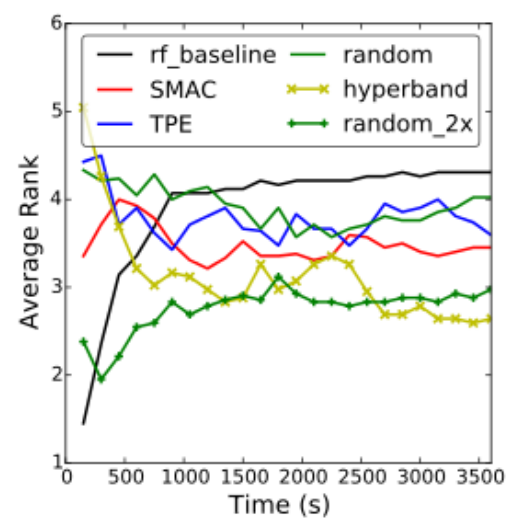
(a)



(b)



(c)



(d)



# Within Scikit-learn

- GridSearchCV
- RandomizedSearchCV
- Searching over Pipelines
- Built-in parameter ranges (coming)

# TODO

Clean separation of:

- Model Search Space
  - Pipeline Search Space
  - Optimization Method
  - Meta-Learning
- 
- Exploit prior knowledge better!
  - Usability
  - Runtime consideration

# TODO

Clean separation of:

- Model Search Space
  - Pipeline Search Space
  - Optimization Method
  - Meta-Learning
- 
- Exploit prior knowledge better!
  - Usability
  - Runtime consideration
  - Data subsampling

# Criticism

Randomized Search works well

Do we need 100 Classifiers?  
Do we need Complex pipelines?

I don't want a black-box!


Making it too easy?



# Material

- Taking the Human Out of the Loop: A Review of Bayesian Optimization (Shahriari, Swersky, Wang, Adams, de Freitas)
- Random Search for Hyper-Parameter Optimization (Bergstra, Bengio)
- Efficient and Robust Automated Machine Learning (Feurer et al) [autosklearn]
- <http://automl.github.io/auto-sklearn/stable/>
- Efficient Hyperparameter Optimization and Infinitely Many Armed Bandits (Lie et. al) [hyperband] <https://arxiv.org/abs/1603.06560>
- Scalable Bayesian Optimization Using Deep Neural Networks [Snoek et al]
- <https://github.com/iaroslav-ai/scikit-optimize-benchmarks>

O'REILLY®



# Introduction to Machine Learning with Python

A GUIDE FOR DATA SCIENTISTS

Andreas C. Müller & Sarah Guido




[amueller.github.io](https://github.com/amueller)



[@amuellerm1](https://twitter.com/amuellerm1)



[@amueller](https://github.com/amueller)



[t3kcit@gmail.com](mailto:t3kcit@gmail.com)

[https://github.com/amueller/talks\\_odt/](https://github.com/amueller/talks_odt/)