

An update on scikit-learn


0.20 and beyond



Andreas Müller
Columbia University, scikit-learn



What is scikit-learn?



Basic API review

Representing Data

one sample

$X =$

| | | | | |
|-----|-----|-----|-----|-----|
| 1.1 | 2.2 | 3.4 | 5.6 | 1.0 |
| 6.7 | 0.5 | 0.4 | 2.6 | 1.6 |
| 2.4 | 9.3 | 7.3 | 6.4 | 2.8 |
| 1.5 | 0.0 | 4.3 | 8.3 | 3.4 |
| 0.5 | 3.5 | 8.1 | 3.6 | 4.6 |
| 5.1 | 9.7 | 3.5 | 7.9 | 5.1 |
| 3.7 | 7.8 | 2.6 | 3.2 | 6.3 |

one feature

$y =$

| |
|-----|
| 1.6 |
| 2.7 |
| 4.4 |
| 0.5 |
| 0.2 |
| 5.6 |
| 6.7 |


outputs / labels

```
clf = RandomForestClassifier()
```

```
clf.fit(X_train, y_train)
```

```
y_pred = clf.predict(X_test)
```

```
clf.score(X_test, y_test)
```



Basic API

`estimator.fit(X, [y])`

`estimator.predict`

`estimator.transform`

Classification

Preprocessing

Regression

Dimensionality reduction

Clustering

Feature selection

Feature extraction

Cross -Validated Grid Search

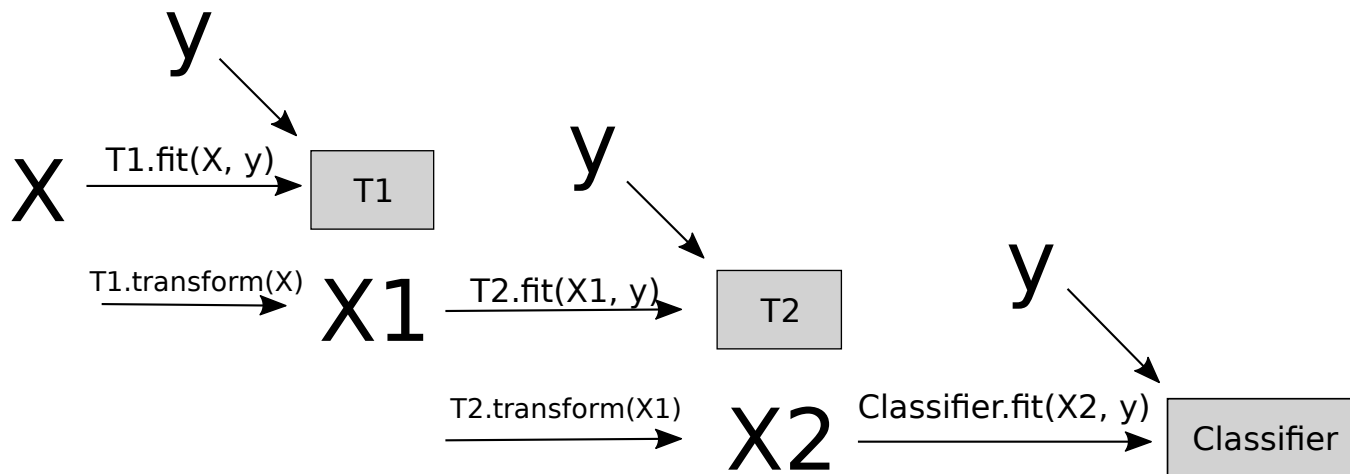
```
from sklearn.grid_search import GridSearchCV
from sklearn.cross_validation import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y)
param_grid = {'C': 10. ** np.arange(-3, 3),
              'gamma': 10. ** np.arange(-3, 3)}
grid = GridSearchCV(SVC(), param_grid=param_grid)
grid.fit(X_train, y_train)
grid.predict(X_test)
grid.score(X_test, y_test)
```


Pipelines


```
pipe = make_pipeline(T1(), T2(), Classifier())
```



```
pipe.fit(X, y)
```



```
pipe.predict(X')
```



Pipelines

```
from sklearn.pipeline import make_pipeline
```

```
pipe = make_pipeline(StandardScaler(), SVC())  
pipe.fit(X_train, y_train)  
pipe.predict(X_test)
```

```
param_grid = {'svc__C': 10. ** np.arange(-3, 3),  
              'svc__gamma': 10. ** np.arange(-3, 3)}
```

```
scaler_pipe = make_pipeline(StandardScaler(), SVC())  
grid = GridSearchCV(scaler_pipe, param_grid=param_grid, cv=5)  
grid.fit(X_train, y_train)
```



```
conda install scikit-learn=0.20rc1 -c conda-forge/label/rc -c conda-forge
```

```
pip install --pre scikit-learn
```

OneHotEncoder for Strings

```
>>> from sklearn.preprocessing import OneHotEncoder
>>> enc = OneHotEncoder(handle_unknown='ignore')
>>> X = [['Male', 1], ['Female', 3], ['Female', 2]]
>>> enc.fit(X)
```

```
>>> enc.transform([['Female', 1], ['Male', 4]]).toarray()
array([[1., 0., 1., 0., 0.],
       [0., 1., 0., 0., 0.]])
>>> enc.inverse_transform([[0, 1, 1, 0, 0], [0, 0, 0, 1, 0]])
array([['Male', 1],
       [None, 2]], dtype=object)
```

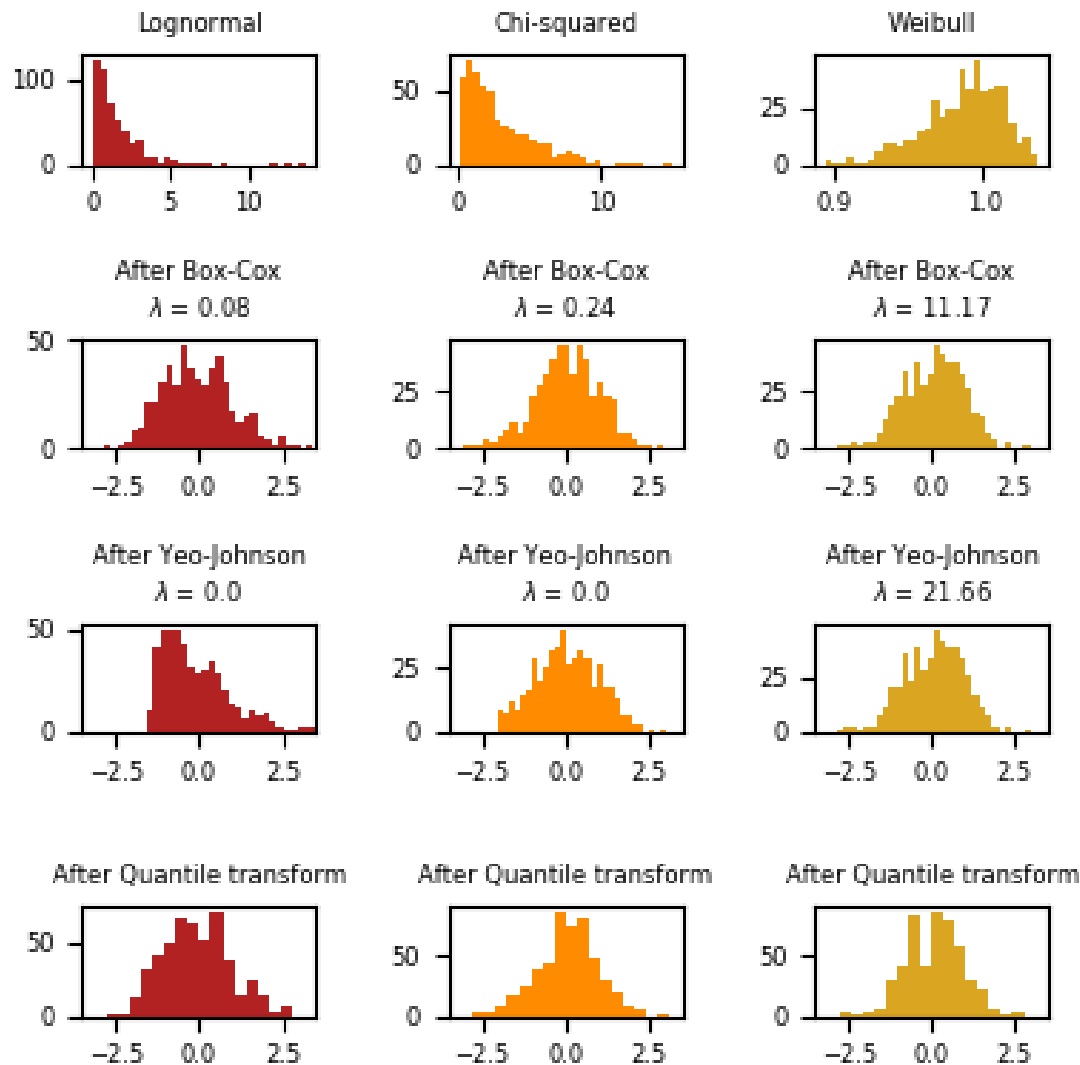
ColumnTransformer

```
numeric_features = ['age', 'fare']
numeric_transformer = make_pipeline(
    SimpleImputer(strategy='median'),
    StandardScaler())

categorical_features = ['embarked', 'sex', 'pclass']
categorical_transformer = make_pipeline(
    SimpleImputer(strategy='constant'),
    OneHotEncoder(sparse=False))

preprocessor = make_column_transformer(
    (numeric_transformer, numeric_features),
    (categorical_transformer, categorical_features),
    remainder='drop')
```

PowerTransformer



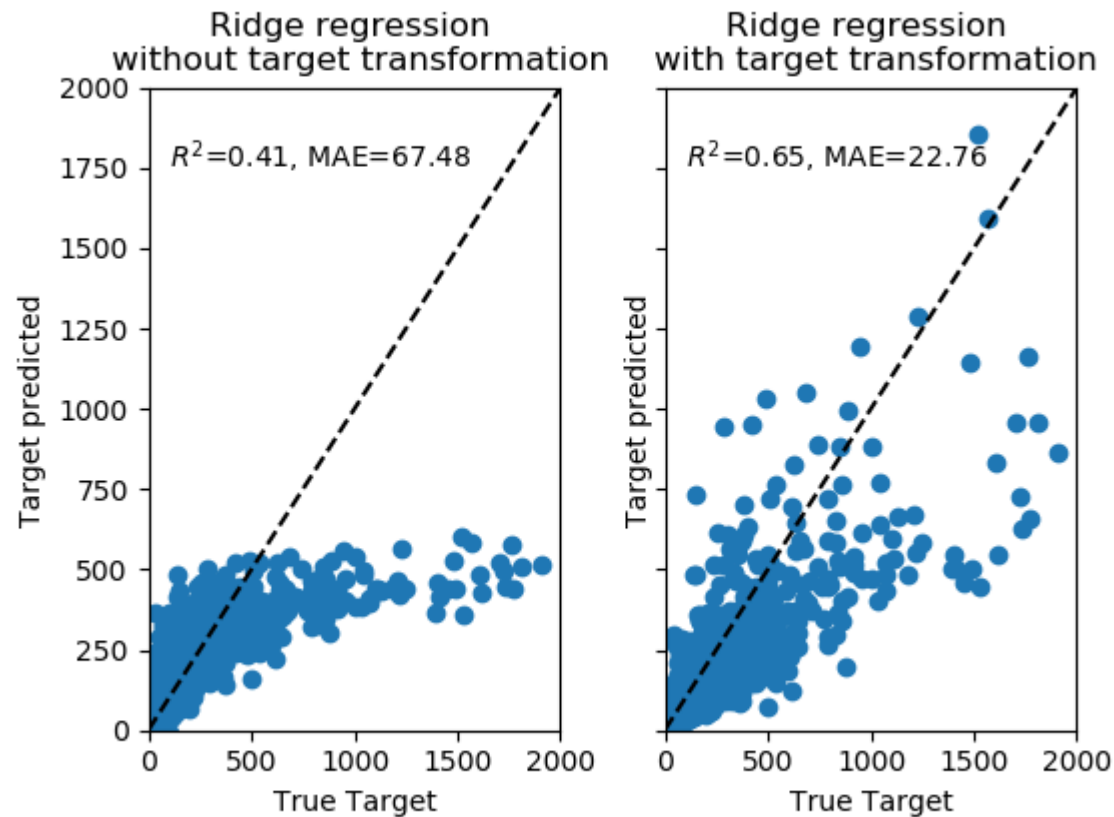
$$y_i^{(\lambda)} = \begin{cases} \frac{y_i^\lambda - 1}{\lambda} & \text{if } \lambda \neq 0, \\ \ln(y_i) & \text{if } \lambda = 0, \end{cases}$$

Missing Value treatment

- Scalers
- SimpleImputer
- MissingIndicator


TransformedTargetRegressor

```
regr_trans = TransformedTargetRegressor(regressor=RidgeCV(),  
                                         func=np.log1p,  
                                         inverse_func=np.expm1)  
  
regr_trans.fit(X_train, y_train)  
y_pred = regr_trans.predict(X_test)
```



Synthetic data

OpenML Dataset Loader



Machine learning, better, together

19929
data sets

Find or add **data** to analyse

46723
tasks

Download or create scientific
tasks

5061
flows

Find or add data analysis **flows**

8442991
runs

Upload and explore all **results**
online.

```
>>> from sklearn.datasets import fetch_openml
>>> mice = fetch_openml('miceprotein', version=4, data_home=custom_data_home)
```

Loky

A Robust and reusable Executor

<https://loky.readthedocs.io/en/stable/>

- An alternative for multiprocessing.pool.Pool and concurrent.futures.ProcessPoolExecutor
- No need for `if __name__ == "__main__":` in scripts
- Deadlock free implementation
- Consistent spawn behavior
- No random crashes with odd BLAS / OpenMP libraries

Global config and working memory

```
>>> import sklearn
>>> with sklearn.config_context(working_memory=128):
...     pass  # do chunked work here
```

sklearn.set_config

sklearn. **set_config** (*assume_finite=None, working_memory=None*)

[\[source\]](#)

Set global scikit-learn configuration

Parameters: **assume_finite** : *bool, optional*

If True, validation for finiteness will be skipped, saving time, but leading to potential crashes. If False, validation for finiteness will be performed, avoiding error. Global default: False.



working_memory : *int, optional*

If set, scikit-learn will attempt to limit the size of temporary arrays to this number of MiB (per job when parallelised), often saving both computation time and memory on expensive operations that can be performed in chunks. Global default: 1024.

Gradient Boosting Early Stopping

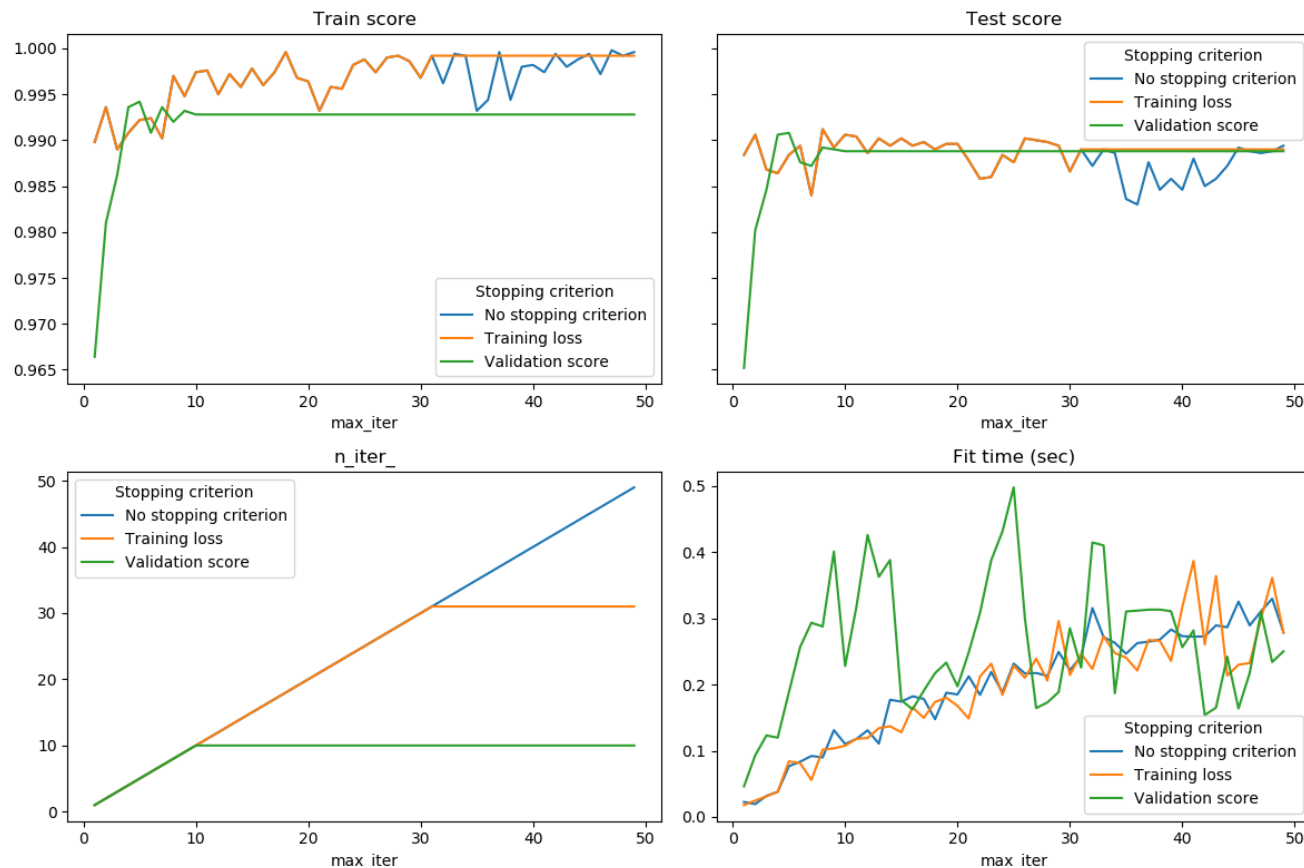
```
GradientBoostingClassifier(n_estimators=n_estimators,  
                             validation_fraction=0.2,  
                             n_iter_no_change=5, tol=0.01,  
                             random_state=0)
```

Activated by “n_iter_no_change”



SGD et al early stopping

```
'Training loss':  
linear_model.SGDClassifier(early_stopping=False, n_iter_no_change=3,  
                             tol=0.1),  
  
'Validation score':  
linear_model.SGDClassifier(early_stopping=True, n_iter_no_change=3,  
                             tol=0.0001, validation_fraction=0.2)
```



[Home](#)[Installation](#)[Documentation](#)[Examples](#)[Fork me on GitHub](#)[Previous](#)[6. Computing ...](#)[Next](#)[Examples](#)[scikit-learn v0.20rc1](#)[Other versions](#)

Please **cite us** if you use
the software.

Glossary of Common Terms and API Elements

[General Concepts](#)[Class APIs and Estimator Types](#)[Target Types](#)[Methods](#)[Parameters](#)[Attributes](#)[Data and sample properties](#)

Glossary of Common Terms and API Elements

« This glossary hopes to definitively represent the tacit and explicit conventions applied in Scikit-learn and its API, while providing a reference for users and contributors. It aims to describe the concepts and either detail their corresponding API or link to other relevant parts of the documentation which do so. By linking to glossary entries from the API Reference and User Guide, we may minimize redundancy and inconsistency.

We begin by listing general concepts (and any that didn't fit elsewhere), but more specific sets of related terms are listed below: [Class APIs and Estimator Types](#), [Target Types](#), [Methods](#), [Parameters](#), [Attributes](#), [Data](#) and [sample properties](#).

General Concepts

1d

1d array

One-dimensional array. A NumPy array whose `.shape` has length 1. A vector.

Better defaults

All random forests:

n_estimators from 10 to 100 (in 0.22)

Cross-validation:

cv from 3 to 5 (in 0.22)

Grid-Search:


iid to False (in 0.22)

Remove iid (in 0.24)

LogisticRegression defaults

- `solver='lbfgs'` (from 'liblinear')
- `multiclass='auto'` (from 'ovr')

“Fixed” iris

4  sklearn/datasets/data/iris.csv



@@ -33,10 +33,10 @@

33 5.4,3.4,1.5,0.4,0

34 5.2,4.1,1.5,0.1,0

35 5.5,4.2,1.4,0.2,0

36 - 4.9,3.1,1.5,0.1,0

37 5.0,3.2,1.2,0.2,0

38 5.5,3.5,1.3,0.2,0

39 - 4.9,3.1,1.5,0.1,0

40 4.4,3.0,1.3,0.2,0

41 5.1,3.4,1.5,0.2,0

42 5.0,3.5,1.3,0.3,0

33 5.4,3.4,1.5,0.4,0

34 5.2,4.1,1.5,0.1,0

35 5.5,4.2,1.4,0.2,0

36 + 4.9,3.1,1.5,0.2,0

37 5.0,3.2,1.2,0.2,0

38 5.5,3.5,1.3,0.2,0

39 + 4.9,3.6,1.4,0.1,0

40 4.4,3.0,1.3,0.2,0

41 5.1,3.4,1.5,0.2,0

42 5.0,3.5,1.3,0.3,0



A note on deprecations ...

[WIP]

Work In Progress



Open

A (draft) Roadmap

<https://github.com/scikit-learn/scikit-learn/wiki/Draft-Roadmap-2018>

Dropping 2.7 (and 3.4!)

Loky / OpenMP

Imbalanced-learn integration

Sampler: To resample a data sets, each sampler implements:

```
data_resampled, targets_resampled = obj.sample(data, targets)
```


Fitting and sampling can also be done in one step:

```
data_resampled, targets_resampled = obj.fit_sample(data, targets)
```

Pandas & Feature Names

```
preprocessor = make_column_transformer(  
    (numeric_transformer, numeric_features),  
    (categorical_transformer, categorical_features),  
    remainder='drop')  
  
preprocessor.fit(X_train)  
preprocessor.get_feature_names()
```


O'REILLY®



Introduction to Machine Learning with Python

A GUIDE FOR DATA SCIENTISTS

Andreas C. Müller & Sarah Guido




amueller.github.io



[@amuellerml](https://twitter.com/amuellerml)



[@amueller](https://github.com/amueller)



[t3kcit
@gmail.com](mailto:t3kcit@gmail.com)