

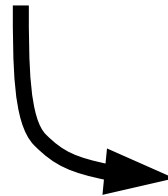
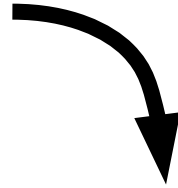





Building scikit-learn

Andreas Müller
Columbia University,
scikit-learn







Classification
Regression
Clustering
Semi-Supervised Learning
Feature Selection
Feature Extraction
Manifold Learning
Dimensionality Reduction
Kernel Approximation
Hyperparameter Optimization
Evaluation Metrics
Out-of-core learning

.....





Alexandre Gramfort

agramfort



Alexander Fabisch

AlexanderFabisch



Alexandre Passos

alextp



Andreas Mueller

amueller



Arnaud Joly

arjoly



Brian Holt

bdholt1



bthirion

bthirion



Chris Filo Gorgolewski

chrisfilo



David Courneau

cournape



Duchesnay

duchesnay



David Warde-Farley

dwf



Fabian Pedregosa

fabianp



Gael Varoquaux

GaelVaroquaux



Gilles Louppe

glouppe



Jake Vanderplas

jakevdp



Jaques Grobler

jaquesgrobler



Jan Hendrik Metzen

jmetzen



Jacob Schreiber

jmschrei



Joel Nothman

jnothman



Kyle Kastner

kastnerkyle



Lars

larsmans



Loïc Estève

lesteve



Shiqiao Du

lucidfrontier45



Mathieu Blondel

mbondel



Manoj Kumar

MechCoder



Noel Dawe

ndawe



Nelle Varoquaux

NelleV



Olivier Grisel

ogrisel



Paolo Losi

paolo-losi



Peter Prettenhofer

pprett



(Venkat) Raghav (Rajagopalan)

raghavrv



Robert Layton

robertlayton



Ron Weiss

ronw



Satrajit Ghosh

satra



sklearn-ci



sklearn-wheels



Tom Dupré la Tour

TomDLT



Vlad Niculae

vene



Virgile Fritsch

VirgileFritsch



Vincent Michel

vmichel



Wei Li

weilinear



Yaroslav Halchenko

yarikoptic

Mission

Commoditize and Democratize Machine Learning

Basic API

Representing Data



one sample

$$X = \begin{pmatrix} 1.1 & 2.2 & 3.4 & 5.6 & 1.0 \\ 6.7 & 0.5 & 0.4 & 2.6 & 1.6 \\ 2.4 & 9.3 & 7.3 & 6.4 & 2.8 \\ 1.5 & 0.0 & 4.3 & 8.3 & 3.4 \\ 0.5 & 3.5 & 8.1 & 3.6 & 4.6 \\ 5.1 & 9.7 & 3.5 & 7.9 & 5.1 \\ 3.7 & 7.8 & 2.6 & 3.2 & 6.3 \end{pmatrix}$$

one feature

$$y = \begin{pmatrix} 1.6 \\ 2.7 \\ 4.4 \\ 0.5 \\ 0.2 \\ 5.6 \\ 6.7 \end{pmatrix}$$


outputs / labels


```
clf = RandomForestClassifier()
```

```
clf.fit(X_train, y_train)
```

```
y_pred = clf.predict(X_test)
```

```
clf.score(X_test, y_test)
```




Unsupervised Transformations

```
pca = PCA()
```

```
pca.fit(X_train)
```

```
X_new = pca.transform(X_test)
```



Core API Summary

`estimator.fit(X, [y])`

`estimator.predict`

`estimator.transform`

Classification

Preprocessing

Regression

Dimensionality reduction

Clustering

Feature selection

Feature extraction

Guiding Ideas

Goals:

Maintainability
Ease of use

Simple things should be simple,
complex things should be possible.

Alan Kay



Three way documentation

1.9. Ensemble methods

The goal of **ensemble methods** is to combine the predictions of several base estimators built with a given learning algorithm in order to improve generalizability / robustness over a single estimator.

Two families of ensemble methods are usually distinguished:

- In **averaging methods**, the driving principle is to build several estimators independently and then to average their predictions. On average, the combined estimator is usually better than any of the single base estimator because its variance is reduced.

Examples: *Bagging methods, Forests of randomized trees, ...*

- By contrast, in **boosting methods**, base estimators are built sequentially and one tries to reduce the bias of the combined estimator. The motivation is to combine several weak models to produce a powerful ensemble.

Examples: *AdaBoost, Gradient Tree Boosting, ...*

`sklearn.ensemble.RandomForestClassifier`

```
class sklearn.ensemble. RandomForestClassifier (n_estimators=10, criterion='gini',
max_depth=None, min_samples_split=2, min_samples_leaf=1,
min_weight_fraction_leaf=0.0, max_features='auto', max_leaf_nodes=None,
bootstrap=True, oob_score=False, n_jobs=1, random_state=None, verbose=0,
warm_start=False)
```

[\[source\]](#)

A random forest classifier.

A random forest is a meta estimator that fits a number of decision tree classifiers on various sub-samples of the dataset and use averaging to improve the predictive accuracy and control over-fitting.

Parameters: `n_estimators` : integer, optional (default=10)

The number of trees in the forest.


criterion : string, optional (default="gini")

The function to measure the quality of a split. Supported criteria are "gini" for the Gini impurity and "entropy" for the information gain. Note: this parameter is tree-specific.

Examples




*Recognizing
hand-written digits*



*Plot classification
probability*



Classifier comparison



*Linear and Quadratic
Discriminant Analysis
with confidence
ellipsoid*

Ease of Installation

Simplicity

```
lr = LogisticRegression()  
lr.fit(X_train, y_train)  
lr.score(X_test, y_test)
```

Consistency

```
grid = GridSearchCV(svm, param_grid)
grid.fit(X_train, y_train)
grid.score(X_test, y_test)
```

Composition

```
feature_selection = RFECV(NaiveBayes())  
pipe = make_pipeline(feature_selection,  
                      RandomForestClassifier())  
grid = GridSearchCV(pipe, param_grid)
```

Default Parameters

```
In [2]: clf = SVC()  
        clf.fit(X_train, y_train)
```

```
SVC(self, C=1.0, kernel='rbf', degree=3, gamma=0.0, coef0=0.0,  
     shrinking=True, probability=False, tol=0.001, cache_size=200,  
     class_weight=None, verbose=False, max_iter=-1, random_state=None)
```

Flat Class Hierarchy, Few Types

- Numpy arrays / sparse matrices
- Estimators
- [Cross-validation objects]
- [Scorers]



Programs should be written for people to read,
and only incidentally for machines to execute.

Harold Abelson

Maintainability


THE LIFE OF A SOFTWARE ENGINEER.

CLEAN SLATE. SOLID FOUNDATIONS. THIS TIME I WILL BUILD THINGS THE RIGHT WAY.




MUCH LATER...

OH MY. I'VE DONE IT AGAIN, HAVEN'T I ?



Feature Creep

Scoping



Testing & Continuous Integration

Add more commits by pushing to the `more_repr` branch on `amueller/scikit-learn`.



All checks have passed

3 successful checks

[Hide all checks](#)



ci/circleci — Your tests passed on CircleCI

[Details](#)



continuous-integration/appveyor/pr — AppVeyor build succeeded

[Details](#)



continuous-integration/travis-ci/pr — The Travis CI build passed

[Details](#)



This branch has no conflicts with the base branch

Merging can be performed automatically.

Squash and merge



or view [command line instructions](#).



delete

avoid code;
avoid code rot!

API implementation & contracts

Method Chaining

Fit must return self:

```
X_scaled = StandardScaler().fit(X).transform(X)
```

```
pred = SVC().fit(X_train, y_train).score(  
    X_test, y_test)
```

fit resets

```
from sklearn.tree import DecisionTreeClassifier
tree = DecisionTreeClassifier(max_depth=10)
tree.fit(iris.data, iris.target)
score_iris = tree.score(iris.data, iris.target)
tree.fit(digits.data, digits.target)
score_digits = tree.score(digits.data, digits.target)
```

(only) fit mutates self

- `__init__` only remembers construction parameters
- Transform / score etc don't change object
- Fit returns self but mutates object!

Estimated attributes

Attributes

`components_` : array, shape (n_components, n_features)

Principal axes in feature space, representing the directions of maximum variance in the data. The components are sorted by ``explained_variance``.

`explained_variance_` : array, shape (n_components,)

The amount of variance explained by each of the selected components.

Equal to n_components largest eigenvalues of the covariance matrix of X.

.. versionadded:: 0.18

`explained_variance_ratio_` : array, shape (n_components,)

Percentage of variance explained by each of the selected components.

If ``n_components`` is not set then all components are stored and the sum of the ratios is equal to 1.0.

`singular_values_` : array, shape (n_components,)

The singular values corresponding to each of the selected components. The singular values are equal to the 2-norms of the ``n_components`` variables in the lower-dimensional space.

Leaves two kinds of attributes:
arguments to `__init__`; things estimated during fit

fit_transform / fit_predict

In general: computational shortcut:

```
pca = PCA()  
pca.fit(X)  
X_pca = pca.transform(X)  
X_pca2 = pca.fit_transform(X)
```

Clustering / Manifold learning: Not inductive.

```
tsne = TSNE()  
X_tsne = tsne.fit_transform(X)
```

```
dbscan = DBSCAN()  
cluster_labels = dbscan.fit_predict(X)
```

check_estimator

```
class TemplateClassifier(BaseEstimator, ClassifierMixin):

    def __init__(self, demo_param='demo'):
        self.demo_param = demo_param

    def fit(self, X, y):

        # Check that X and y have correct shape
        X, y = check_X_y(X, y)
        # Store the classes seen during fit
        self.classes_ = unique_labels(y)

        self.X_ = X
        self.y_ = y
        # Return the classifier
        return self

    def predict(self, X):

        closest = np.argmin(euclidean_distances(X, self.X_), axis=1)
        return self.y_[closest]
```

`check_estimator(TemplateClassifier)`

AssertionError: Error message does not include the expected string: 'fit';
Observed error message: "'TemplateClassifier' object has no attribute 'X_'"

Development Practices

Standards for OSS

Everything discussed in the open.
Every convention and process documented.

Development guide

<http://scikit-learn.org/dev/developers/contributing.html>

Contains:

- API details
- Bug report guidelines
- PR guidelines
- Reviewing guidelines
- How to find issues
- Details of CI

Even more at <http://scikit-learn.org/dev/developers/>

Deprecations / backward compatibility

- Don't change any behavior (except bug fixes)


```
@property
@deprecated("Attribute labels_ was deprecated in version 0.13 and "
            "will be removed in 0.15. Use 'classes_' instead")
def labels_(self):
    return self.classes_
```

(my personal) Roadmap

Make simple things simple again!

Missing Values
Categorical Variables
Plotting Tools
Pandas Integration
Feature Names

O'REILLY®



Introduction to Machine Learning with Python

A GUIDE FOR DATA SCIENTISTS

Andreas C. Müller & Sarah Guido




amueller.github.io



[@amuellerm1](https://twitter.com/amuellerm1)



[@amueller](https://github.com/amueller)



t3kcit@gmail.com